

# Bayesian Methods & Multiclass

Mengye Ren

NYU

Oct 31, 2023

# Announcement

- Schedule your project consultation soon.
- Use the provided template! (if your final report fails to use template then there will be marks off)
- Homework 3 is released and due Nov 14 11:59AM.

# Recap

- Bayesian modeling adds a prior on the parameters.
- Models the distribution of parameters.
- Bayes Rule:

$$p(y | x) = \frac{p(x | y)p(y)}{p(x)}$$

- 

$$p(\theta | \mathcal{D}) = \frac{p(\mathcal{D} | \theta)p(\theta)}{p(\mathcal{D})}.$$

- 

$$\underbrace{p(\theta | \mathcal{D})}_{\text{posterior}} \propto \underbrace{p(\mathcal{D} | \theta)}_{\text{likelihood}} \underbrace{p(\theta)}_{\text{prior}}.$$

- Conjugate prior: Having the same form of distribution as the posterior.

# Bayesian Point Estimates

- We have the posterior distribution  $\theta \mid \mathcal{D}$ .
- What if someone asks us for a point estimate  $\hat{\theta}$  for  $\theta$ ?
- Common options:
  - **posterior mean**  $\hat{\theta} = \mathbb{E}[\theta \mid \mathcal{D}]$
  - **maximum a posteriori (MAP) estimate**  $\hat{\theta} = \arg \max_{\theta} p(\theta \mid \mathcal{D})$ 
    - Note: this is the **mode** of the posterior distribution

## What else can we do with a posterior?

- Look at it: display uncertainty estimates to our client
- Extract a **credible set** for  $\theta$  (a Bayesian confidence interval).
  - e.g. Interval  $[a, b]$  is a 95% **credible set** if

$$\mathbb{P}(\theta \in [a, b] \mid \mathcal{D}) \geq 0.95$$

- Select a point estimate using **Bayesian decision theory**:
  - Choose a loss function.
  - Find action **minimizing expected risk w.r.t. posterior**

# Bayesian Decision Theory

---

# Bayesian Decision Theory

- Ingredients:
  - **Parameter space**  $\Theta$ .
  - **Prior**: Distribution  $p(\theta)$  on  $\Theta$ .
  - **Action space**  $\mathcal{A}$ .
  - **Loss function**:  $\ell : \mathcal{A} \times \Theta \rightarrow \mathbb{R}$ .
- The **posterior risk** of an action  $a \in \mathcal{A}$  is

$$\begin{aligned} r(a) &:= \mathbb{E}[\ell(\theta, a) \mid \mathcal{D}] \\ &= \int \ell(\theta, a) p(\theta \mid \mathcal{D}) d\theta. \end{aligned}$$

- It's the **expected loss under the posterior**.
- A **Bayes action**  $a^*$  is an action that minimizes posterior risk:

$$r(a^*) = \min_{a \in \mathcal{A}} r(a)$$

# Bayesian Point Estimation

- General Setup:
  - Data  $\mathcal{D}$  generated by  $p(y | \theta)$ , for unknown  $\theta \in \Theta$ .
  - We want to produce a **point estimate** for  $\theta$ .
- Choose:
  - **Prior**  $p(\theta)$  on  $\Theta = \mathbb{R}$ .
  - **Loss**  $\ell(\hat{\theta}, \theta)$
- Find **action**  $\hat{\theta} \in \Theta$  that minimizes the **posterior risk**:

$$\begin{aligned} r(\hat{\theta}) &= \mathbb{E}[\ell(\hat{\theta}, \theta) | \mathcal{D}] \\ &= \int \ell(\hat{\theta}, \theta) p(\theta | \mathcal{D}) d\theta \end{aligned}$$



# Important Cases

- Squared Loss :  $\ell(\hat{\theta}, \theta) = (\theta - \hat{\theta})^2 \Rightarrow$  posterior mean
- Zero-one Loss:  $\ell(\theta, \hat{\theta}) = \mathbb{1}[\theta \neq \hat{\theta}] \Rightarrow$  posterior mode
- Absolute Loss :  $\ell(\hat{\theta}, \theta) = |\theta - \hat{\theta}| \Rightarrow$  posterior median
- Optimal decision depends on the loss function and the posterior distribution.
- We will derive the square loss case next.

## Bayesian Point Estimation: Square Loss

- Find **action**  $\hat{\theta} \in \Theta$  that minimizes **posterior risk**

$$r(\hat{\theta}) = \int (\theta - \hat{\theta})^2 p(\theta | \mathcal{D}) d\theta.$$

- Differentiate:

$$\begin{aligned} \frac{dr(\hat{\theta})}{d\hat{\theta}} &= - \int 2(\theta - \hat{\theta}) p(\theta | \mathcal{D}) d\theta \\ &= -2 \int \theta p(\theta | \mathcal{D}) d\theta + 2\hat{\theta} \underbrace{\int p(\theta | \mathcal{D}) d\theta}_{=1} \\ &= -2 \int \theta p(\theta | \mathcal{D}) d\theta + 2\hat{\theta} \end{aligned}$$

# Bayesian Point Estimation: Square Loss

- Derivative of posterior risk is

$$\frac{dr(\hat{\theta})}{d\hat{\theta}} = -2 \int \theta p(\theta | \mathcal{D}) d\theta + 2\hat{\theta}.$$

- First order condition  $\frac{dr(\hat{\theta})}{d\hat{\theta}} = 0$  gives

$$\begin{aligned}\hat{\theta} &= \int \theta p(\theta | \mathcal{D}) d\theta \\ &= \mathbb{E}[\theta | \mathcal{D}]\end{aligned}$$

- The **Bayes action** for **square loss** is the posterior mean.

## Interim summary

---

## Recap and Interpretation

- The prior represents belief about  $\theta$  before observing data  $\mathcal{D}$ .
- The posterior represents **rationally updated beliefs** after seeing  $\mathcal{D}$ .
- All inferences and action-taking are based on the posterior distribution.
- In the Bayesian approach,
  - No issue of justifying an estimator.
  - Only choices are
    - **family of distributions**, indexed by  $\Theta$ , and
    - **prior distribution** on  $\Theta$
  - For decision making, we need a **loss function**.

## Recap: Conditional Probability Models

---

# Conditional Probability Modeling

- Input space  $\mathcal{X}$
- Outcome space  $\mathcal{Y}$
- Action space  $\mathcal{A} = \{p(y) \mid p \text{ is a probability distribution on } \mathcal{Y}\}$ .
- Hypothesis space  $\mathcal{F}$  contains prediction functions  $f : \mathcal{X} \rightarrow \mathcal{A}$ .
- Prediction function  $f \in \mathcal{F}$  takes input  $x \in \mathcal{X}$  and produces a **distribution** on  $\mathcal{Y}$
- A **parametric family of conditional densities** is a set

$$\{p(y \mid x, \theta) : \theta \in \Theta\},$$

- where  $p(y \mid x, \theta)$  is a density on **outcome space**  $\mathcal{Y}$  for each  $x$  in **input space**  $\mathcal{X}$ , and
  - $\theta$  is a **parameter** in a [finite dimensional] **parameter space**  $\Theta$ .
- This is the common starting point for either classical or Bayesian regression.

# Classical treatment: Likelihood Function

- **Data:**  $\mathcal{D} = (y_1, \dots, y_n)$
- The probability density for our data  $\mathcal{D}$  is

$$p(\mathcal{D} \mid x_1, \dots, x_n, \theta) = \prod_{i=1}^n p(y_i \mid x_i, \theta).$$

- For fixed  $\mathcal{D}$ , the function  $\theta \mapsto p(\mathcal{D} \mid x, \theta)$  is the **likelihood function**:

$$L_{\mathcal{D}}(\theta) = p(\mathcal{D} \mid x, \theta),$$

where  $x = (x_1, \dots, x_n)$ .



# Maximum Likelihood Estimator

- The **maximum likelihood estimator (MLE)** for  $\theta$  in the family  $\{p(y | x, \theta) | \theta \in \Theta\}$  is

$$\hat{\theta}_{\text{MLE}} = \arg \max_{\theta \in \Theta} L_{\mathcal{D}}(\theta).$$

- MLE corresponds to ERM, if we set the loss to be the negative log-likelihood.
- The corresponding prediction function is

$$\hat{f}(x) = p(y | x, \hat{\theta}_{\text{MLE}}).$$

# Bayesian Conditional Probability Models

---

# Bayesian Conditional Models

- Input space  $\mathcal{X} = \mathbb{R}^d$       Outcome space  $\mathcal{Y} = \mathbb{R}$
- The Bayesian conditional model has two components:
  - A **parametric family of conditional densities**:

$$\{p(y | x, \theta) : \theta \in \Theta\}$$

- A **prior distribution**  $p(\theta)$  on  $\theta \in \Theta$ .

# The Posterior Distribution

- The **prior distribution**  $p(\theta)$  represents our beliefs about  $\theta$  before seeing  $\mathcal{D}$ .
- The **posterior distribution** for  $\theta$  is

$$\begin{aligned} p(\theta \mid \mathcal{D}, x) &\propto p(\mathcal{D} \mid \theta, x) p(\theta) \\ &= \underbrace{L_{\mathcal{D}}(\theta)}_{\text{likelihood}} \underbrace{p(\theta)}_{\text{prior}} \end{aligned}$$

- Posterior represents the **rationally updated beliefs** after seeing  $\mathcal{D}$ .
- Each  $\theta$  corresponds to a prediction function,
  - i.e. the conditional distribution function  $p(y \mid x, \theta)$ .

# Point Estimates of Parameter

- What if we want point estimates of  $\theta$ ?
- We can use **Bayesian decision theory** to derive point estimates.
- We may want to use
  - $\hat{\theta} = \mathbb{E}[\theta \mid \mathcal{D}, x]$  (the posterior mean estimate)
  - $\hat{\theta} = \text{median}[\theta \mid \mathcal{D}, x]$
  - $\hat{\theta} = \arg \max_{\theta \in \Theta} p(\theta \mid \mathcal{D}, x)$  (the MAP estimate)
- depending on our loss function.

## Back to the basic question - Bayesian Prediction Function

- Find a function takes input  $x \in \mathcal{X}$  and produces a **distribution** on  $\mathcal{Y}$
- In the frequentist approach:
  - Choose family of conditional probability densities (hypothesis space).
  - Select one conditional probability from family, e.g. using MLE.
- In the Bayesian setting:

- We choose a parametric family of conditional densities

$$\{p(y | x, \theta) : \theta \in \Theta\},$$

- and a prior distribution  $p(\theta)$  on this set.
- Having set our Bayesian model, how do we predict a distribution on  $y$  for input  $x$ ?
- We don't need to make a discrete selection from the hypothesis space: we **maintain uncertainty**.

# The Prior Predictive Distribution

- Suppose we have not yet observed any data.
- In the Bayesian setting, we can still produce a prediction function.
- The **prior predictive distribution** is given by

$$x \mapsto p(y | x) = \int p(y | x; \theta) p(\theta) d\theta.$$

- This is an average of all conditional densities in our family, weighted by the prior.

# The Posterior Predictive Distribution

- Suppose we've already seen data  $\mathcal{D}$ .
- The **posterior predictive distribution** is given by

$$x \mapsto p(y \mid x, \mathcal{D}) = \int p(y \mid x; \theta) p(\theta \mid \mathcal{D}) d\theta.$$

- This is an average of all conditional densities in our family, weighted by the posterior.



## Comparison to Frequentist Approach

- In Bayesian statistics we have two distributions on  $\Theta$ :
  - the prior distribution  $p(\theta)$
  - the posterior distribution  $p(\theta | \mathcal{D})$ .
- These distributions over parameters correspond to distributions on the hypothesis space:

$$\{p(y | x, \theta) : \theta \in \Theta\}.$$

- In the frequentist approach, we choose  $\hat{\theta} \in \Theta$ , and predict

$$p(y | x, \hat{\theta}(\mathcal{D})).$$

- In the Bayesian approach, we integrate out over  $\Theta$  w.r.t.  $p(\theta | \mathcal{D})$  and predict with

$$p(y | x, \mathcal{D}) = \int p(y | x; \theta) p(\theta | \mathcal{D}) d\theta$$

## What if we don't want a full distribution on $y$ ?

- Once we have a predictive distribution  $p(y | x, \mathcal{D})$ ,
  - we can easily generate single point predictions.
- $x \mapsto \mathbb{E}[y | x, \mathcal{D}]$ , to minimize expected square error.
- $x \mapsto \text{median}[y | x, \mathcal{D}]$ , to minimize expected absolute error
- $x \mapsto \arg \max_{y \in \mathcal{Y}} p(y | x, \mathcal{D})$ , to minimize expected 0/1 loss
- Each of these can be derived from  $p(y | x, \mathcal{D})$ .

## Gaussian Regression Example

---

## Example in 1-Dimension: Setup

- Input space  $\mathcal{X} = [-1, 1]$       Output space  $\mathcal{Y} = \mathbb{R}$
- Given  $x$ , the world generates  $y$  as

$$y = w_0 + w_1 x + \varepsilon,$$

where  $\varepsilon \sim \mathcal{N}(0, 0.2^2)$ .


- Written another way, the **conditional probability model** is

$$y \mid x, w_0, w_1 \sim \mathcal{N}(w_0 + w_1 x, 0.2^2).$$

- What's the parameter space?  $\mathbb{R}^2$ .
- **Prior distribution:**  $w = (w_0, w_1) \sim \mathcal{N}(0, \frac{1}{2}I)$


## Example in 1-Dimension: Prior Situation

- **Prior distribution:**  $w = (w_0, w_1) \sim \mathcal{N}(0, \frac{1}{2}I)$  (Illustrated on left)




- On right,  $y(x) = \mathbb{E}[y \mid x, w] = w_0 + w_1 x$ , for randomly chosen  $w \sim p(w) = \mathcal{N}(0, \frac{1}{2}I)$ .

## Example in 1-Dimension: 1 Observation



- On left: posterior distribution; white cross indicates true parameters
- On right:
  - blue circle indicates the training observation
  - red lines,  $y(x) = \mathbb{E}[y | x, w] = w_0 + w_1 x$ , for randomly chosen  $w \sim p(w|\mathcal{D})$  (posterior)

## Example in 1-Dimension: 2 and 20 Observations



Bishop's PRML Fig 3.7

## Gaussian Regression: Closed form

---



# Closed Form for Posterior

- Model:

$$\begin{aligned} w &\sim \mathcal{N}(0, \Sigma_0) \\ y_i | x, w &\text{ i.i.d. } \mathcal{N}(w^T x_i, \sigma^2) \end{aligned}$$

- Design matrix  $X$       Response column vector  $y$
- **Posterior distribution is a Gaussian distribution:**

$$\begin{aligned} w | \mathcal{D} &\sim \mathcal{N}(\mu_P, \Sigma_P) \\ \mu_P &= (X^T X + \sigma^2 \Sigma_0^{-1})^{-1} X^T y \\ \Sigma_P &= (\sigma^{-2} X^T X + \Sigma_0^{-1})^{-1} \end{aligned}$$

- **Posterior Variance  $\Sigma_P$  gives us a natural uncertainty measure.**

# Closed Form for Posterior

- Posterior distribution is a **Gaussian distribution**:

$$w | \mathcal{D} \sim \mathcal{N}(\mu_P, \Sigma_P)$$

$$\mu_P = (X^T X + \sigma^2 \Sigma_0^{-1})^{-1} X^T y$$

$$\Sigma_P = (\sigma^{-2} X^T X + \Sigma_0^{-1})^{-1}$$

- If we want point estimates of  $w$ , **MAP estimator** and the **posterior mean** are given by

$$\hat{w} = \mu_P = (X^T X + \sigma^2 \Sigma_0^{-1})^{-1} X^T y$$

- For the prior variance  $\Sigma_0 = \frac{\sigma^2}{\lambda} I$ , we get

$$\hat{w} = \mu_P = (X^T X + \lambda I)^{-1} X^T y,$$

which is of course the ridge regression solution.

## Connection the MAP to Ridge Regression

- The **Posterior density** on  $w$  for  $\Sigma_0 = \frac{\sigma^2}{\lambda} I$ :

$$p(w | \mathcal{D}) \propto \underbrace{\exp\left(-\frac{\lambda}{2\sigma^2} \|w\|^2\right)}_{\text{prior}} \underbrace{\prod_{i=1}^n \exp\left(-\frac{(y_i - w^T x_i)^2}{2\sigma^2}\right)}_{\text{likelihood}}$$

- To find the **MAP**, we minimize the negative log posterior:

$$\begin{aligned}\hat{w}_{\text{MAP}} &= \arg \min_{w \in \mathbb{R}^d} [-\log p(w | \mathcal{D})] \\ &= \arg \min_{w \in \mathbb{R}^d} \underbrace{\sum_{i=1}^n (y_i - w^T x_i)^2}_{\text{log-likelihood}} + \underbrace{\lambda \|w\|^2}_{\text{log-prior}}\end{aligned}$$

- Which is the ridge regression objective.

# Predictive Posterior Distribution

- Given a new input point  $x_{\text{new}}$ , how do we predict  $y_{\text{new}}$  ?
- **Predictive distribution**

$$\begin{aligned} p(y_{\text{new}} | x_{\text{new}}, \mathcal{D}) &= \int p(y_{\text{new}} | x_{\text{new}}, w, \mathcal{D}) p(w | \mathcal{D}) dw \\ &= \int p(y_{\text{new}} | x_{\text{new}}, w) p(w | \mathcal{D}) dw \end{aligned}$$

- For Gaussian regression, predictive distribution has closed form.

# Closed Form for Predictive Distribution

- **Model:**

$$\begin{aligned} w &\sim \mathcal{N}(0, \Sigma_0) \\ y_i | x, w &\text{ i.i.d. } \mathcal{N}(w^T x_i, \sigma^2) \end{aligned}$$

- **Predictive Distribution**

$$p(y_{\text{new}} | x_{\text{new}}, \mathcal{D}) = \int p(y_{\text{new}} | x_{\text{new}}, w) p(w | \mathcal{D}) dw.$$


- Averages over prediction for each  $w$ , weighted by posterior distribution.

- **Closed form:**

$$\begin{aligned} y_{\text{new}} | x_{\text{new}}, \mathcal{D} &\sim \mathcal{N}(\eta_{\text{new}}, \sigma_{\text{new}}^2) \\ \eta_{\text{new}} &= \mu_P^T x_{\text{new}} \\ \sigma_{\text{new}}^2 &= \underbrace{x_{\text{new}}^T \Sigma_P x_{\text{new}}}_{\text{from variance in } w} + \underbrace{\sigma^2}_{\text{inherent variance in } y} \end{aligned}$$

# Bayesian Regression Provides Uncertainty Estimates

- With predictive distributions, we can give mean prediction with error bands:



## Multi-class Overview

- So far, most algorithms we've learned are designed for binary classification.
  - Sentiment analysis (positive vs. negative)
  - Spam filter (spam vs. non-spam)
- Many real-world problems have more than two classes.
  - Document classification (over 10 classes)
  - Object recognition (over 20k classes)
  - Face recognition (millions of classes)
- What are some potential issues when we have a large number of classes?
  - Computation cost
  - Class imbalance
  - Different cost of errors



# Today's lecture

- How to *reduce* multiclass classification to binary classification?
  - We can think of binary classifier or linear regression as a black box. Naive ways:
  - E.g. multiple binary classifiers produce a binary code for each class (000, 001, 010)
  - E.g. a linear regression produces a numerical value for each class (1.0, 2.0, 3.0)
- How do we *generalize* binary classification algorithm to the multiclass setting?
  - We also need to think about the loss function.
- Example of very large output space: structured prediction.
  - Multi-class: Mutually exclusive class structure.
  - Text: Temporal relational structure.

## Reduction to Binary Classification

---

# One-vs-All / One-vs-Rest

## Setting

- Input space:  $\mathcal{X}$
- Output space:  $\mathcal{Y} = \{1, \dots, k\}$

## Training

- Train  $k$  binary classifiers, one for each class:  $h_1, \dots, h_k : \mathcal{X} \rightarrow \mathbb{R}$ .
- Classifier  $h_i$  distinguishes class  $i$  (+1) from the rest (-1).

## Prediction


- Majority vote:

$$h(x) = \arg \max_{i \in \{1, \dots, k\}} h_i(x)$$

- Ties can be broken arbitrarily.

## OvA: 3-class example (linear classifier)




Consider a dataset with three classes:



**Assumption:** each class is linearly separable from the rest.


Ideal case: only target class has positive score.

Train OvA classifiers:







## OvA: 4-class non linearly separable example

Consider a dataset with four classes:



Cannot separate **red** points from the rest.  
Which classes might have low accuracy?

Train OvA classifiers:



# All vs All / One vs One / All pairs

## Setting

- Input space:  $\mathcal{X}$
- Output space:  $\mathcal{Y} = \{1, \dots, k\}$

## Training

- Train  $\binom{k}{2}$  binary classifiers, one for each pair:  $h_{ij} : \mathcal{X} \rightarrow \mathbb{R}$  for  $i \in [1, k]$  and  $j \in [i+1, k]$ .
- Classifier  $h_{ij}$  distinguishes class  $i$  (+1) from class  $j$  (-1).

## Prediction


- Majority vote (each class gets  $k-1$  votes)

$$h(x) = \arg \max_{i \in \{1, \dots, k\}} \sum_{j \neq i} \underbrace{h_{ij}(x) \mathbb{I}\{i < j\}}_{\text{class } i \text{ is } +1} - \underbrace{h_{ji}(x) \mathbb{I}\{j < i\}}_{\text{class } i \text{ is } -1}$$

- Tournament
- Ties can be broken arbitrarily.


## AvA: four-class example

Consider a dataset with four classes:



**Assumption:** each pair of classes are linearly separable.  
More expressive than OvA.

What's the decision region for the red class?



# OvA vs AvA

		OvA	AvA
computation	train	$O(kB_{\text{train}}(n))$	$O(k^2B_{\text{train}}(n/k))$
	test	$O(kB_{\text{test}})$	$O(k^2B_{\text{test}})$
challenges	train	class imbalance	small training set
	test	calibration / scale tie breaking	

Lack theoretical justification but simple to implement and works well in practice (when # classes is small).



## Code word for labels

Using the reduction approach, can you train fewer than  $k$  binary classifiers?

**Key idea:** Encode labels as binary codes and predict the code bits directly.

OvA encoding:

class	$h_1$	$h_2$	$h_3$	$h_4$
1	1	0	0	0
2	0	1	0	0
3	0	0	1	0
4	0	0	0	1

OvA uses  $k$  bits to encode each label, what's the minimal number of bits you can use?

# Error correcting output codes (ECOC)

Example: 8 classes, 6-bit code

class	$h_1$	$h_2$	$h_3$	$h_4$	$h_5$	$h_6$
1	0	0	0	1	0	0
2	1	0	0	0	0	0
3	0	1	1	0	1	0
4	1	1	0	0	0	0
5	1	1	0	0	1	0
6	0	0	1	1	0	1
7	0	0	1	0	0	0
8	0	1	0	1	0	0

**Training** Binary classifier  $h_i$ :

- +1: classes whose  $i$ -th bit is 1
- -1: classes whose  $i$ -th bit is 0

**Prediction** Closest label in terms of Hamming distance.

$h_1$	$h_2$	$h_3$	$h_4$	$h_5$	$h_6$
0	1	1	0	1	1

**Code design** Want good binary classifiers.

## Error correcting output codes: summary

- Computationally more efficient than OvA (a special case of ECOC). Better for large  $k$ .
- Why not use the minimal number of bits ( $\log_2 k$ )?
  - If the minimum Hamming distance between any pair of code word is  $d$ , then it can correct  $\lfloor \frac{d-1}{2} \rfloor$  errors.
  - In plain words, if rows are far from each other, ECOC is robust to errors.
- Trade-off between code distance and binary classification performance.
- Nice theoretical results [Allwein et al., 2000] (also incorporates AvA).

Reduction-based approaches:

- Reducing multiclass classification to binary classification: OvA, AvA
- Key is to design “natural” binary classification problems without large computation cost.

But,

- Unclear how to generalize to extremely large # of classes.
- ImageNet: >20k labels; Wikipedia: >1M categories.

Next, generalize previous algorithms to multiclass settings.

## Multiclass Loss

# Binary Logistic Regression

- Given an input  $x$ , we would like to output a classification between  $(0,1)$ .

$$f(x) = \textit{sigmoid}(z) = \frac{1}{1 + \exp(-z)} = \frac{1}{1 + \exp(-w^\top x - b)}. \quad (1)$$

- The other class is represented in  $1 - f(x)$ :

$$1 - f(x) = \frac{\exp(-w^\top x - b)}{1 + \exp(-w^\top x - b)} = \frac{1}{1 + \exp(w^\top x + b)} = \textit{sigmoid}(-z). \quad (2)$$

- Another way to view: one class has  $(+w, +b)$  and the other class has  $(-w, -b)$ .

# Multi-class Logistic Regression

- Now what if we have one  $w_c$  for each class  $c$ ?

$$f_c(x) = \frac{\exp(w_c^\top x) + b_c}{\sum_c \exp(w_c^\top x + b_c)} \quad (3)$$

- Also called “softmax” in neural networks.
- Loss function:  $L = \sum_i -y_c^{(i)} \log f_c(x^{(i)})$
- Gradient:  $\frac{\partial L}{\partial z} = f - y$ . Recall: MSE loss.

## Comparison to OvA

- **Base Hypothesis Space:**  $\mathcal{H} = \{h : \mathcal{X} \rightarrow \mathbb{R}\}$  (score functions).
- **Multiclass Hypothesis Space** (for  $k$  classes):

$$\mathcal{F} = \left\{ x \mapsto \arg \max_i h_i(x) \mid h_1, \dots, h_k \in \mathcal{H} \right\}$$

- Intuitively,  $h_i(x)$  scores how likely  $x$  is to be from class  $i$ .
- OvA objective:  $h_i(x) > 0$  for  $x$  with label  $i$  and  $h_i(x) < 0$  for  $x$  with all other labels.
- At test time, to predict  $(x, i)$  correctly we only need

$$h_i(x) > h_j(x) \quad \forall j \neq i. \tag{4}$$



# Multiclass Perceptron

- Base linear predictors:  $h_i(x) = w_i^T x$  ( $w \in \mathbb{R}^d$ ).
- Multiclass perceptron:

Given a multiclass dataset  $\mathcal{D} = \{(x, y)\}$ ;

Initialize  $w \leftarrow 0$ ;

**for**  $iter = 1, 2, \dots, T$  **do**

**for**  $(x, y) \in \mathcal{D}$  **do**

$\hat{y} = \arg \max_{y' \in \mathcal{Y}} w_{y'}^T x$ ;

**if**  $\hat{y} \neq y$  **then** // We've made a mistake

$w_y \leftarrow w_y + x$  ; // Move the target-class scorer towards  $x$

$w_{\hat{y}} \leftarrow w_{\hat{y}} - x$  ; // Move the wrong-class scorer away from  $x$

**end**

**end**

**end**

# Rewrite the scoring function

- Remember that we want to scale to very large # of classes and reuse algorithms and analysis for binary classification
  - $\Rightarrow$  a **single weight vector** is desired
- How to rewrite the equation such that we have one  $w$  instead of  $k$ ?

$$w_i^T x = w^T \psi(x, i) \quad (5)$$

$$h_i(x) = h(x, i) \quad (6)$$

- Encode labels in the feature space.
- Score for each label  $\rightarrow$  score for the “*compatibility*” of a label and an input.

# The Multivector Construction

How to construct the feature map  $\psi$ ?

- What if we stack  $w_i$ 's together (e.g.,  $x \in \mathbb{R}^2, y = \{1, 2, 3\}$ )

$$w = \left( \underbrace{-\frac{\sqrt{2}}{2}, \frac{\sqrt{2}}{2}}_{w_1}, \underbrace{0, 1}_{w_2}, \underbrace{\frac{\sqrt{2}}{2}, \frac{\sqrt{2}}{2}}_{w_3} \right)$$

- And then do the following:  $\Psi: \mathbb{R}^2 \times \{1, 2, 3\} \rightarrow \mathbb{R}^6$  defined by

$$\Psi(x, 1) := (x_1, x_2, 0, 0, 0, 0)$$

$$\Psi(x, 2) := (0, 0, x_1, x_2, 0, 0)$$

$$\Psi(x, 3) := (0, 0, 0, 0, x_1, x_2)$$

- Then  $\langle w, \Psi(x, y) \rangle = \langle w_y, x \rangle$ , which is what we want.

## Rewrite multiclass perceptron

Multiclass perceptron using the multivector construction.

Given a multiclass dataset  $\mathcal{D} = \{(x, y)\}$ ;

Initialize  $w \leftarrow 0$ ;

**for**  $iter = 1, 2, \dots, T$  **do**

**for**  $(x, y) \in \mathcal{D}$  **do**

$\hat{y} = \arg \max_{y' \in \mathcal{Y}} w^T \psi(x, y')$  ; // Equivalent to  $\arg \max_{y' \in \mathcal{Y}} w_{y'}^T x$

**if**  $\hat{y} \neq y$  **then** // We've made a mistake

$w \leftarrow w + \psi(x, y)$  ; // Move the scorer towards  $\psi(x, y)$

$w \leftarrow w - \psi(x, \hat{y})$  ; // Move the scorer away from  $\psi(x, \hat{y})$

**end**

**end**

**end**

**Exercise:** What is the base binary classification problem in multiclass perceptron?

Toy multiclass example: Part-of-speech classification

- $\mathcal{X} = \{\text{All possible words}\}$
- $\mathcal{Y} = \{\text{NOUN, VERB, ADJECTIVE, } \dots\}$ .
- Features of  $x \in \mathcal{X}$ : [The word itself], ENDS\_IN\_ly, ENDS\_IN\_ness, ...

How to construct the feature vector?

- Multivector construction:  $w \in \mathbb{R}^{d \times k}$ —**doesn't scale**.
- Directly design features for each class.

$$\Psi(x, y) = (\psi_1(x, y), \psi_2(x, y), \psi_3(x, y), \dots, \psi_d(x, y)) \quad (7)$$

- Size can be bounded by  $d$ .

# Features

Sample training data:

The boy grabbed the apple and ran away quickly .

Feature:

$$\psi_1(x, y) = \mathbb{1}[x = \text{apple AND } y = \text{NOUN}]$$

$$\psi_2(x, y) = \mathbb{1}[x = \text{run AND } y = \text{NOUN}]$$

$$\psi_3(x, y) = \mathbb{1}[x = \text{run AND } y = \text{VERB}]$$

$$\psi_4(x, y) = \mathbb{1}[x \text{ ENDS\_IN\_ly AND } y = \text{ADVERB}]$$

...

- E.g.,  $\Psi(x = \text{run}, y = \text{NOUN}) = (0, 1, 0, 0, \dots)$
- After training, what's  $w_1, w_2, w_3, w_4$ ?
- No need to include features unseen in training data.

## Feature templates: implementation

- Flexible, e.g., neighboring words, suffix/prefix.
- “Read off” features from the training data.
- Often sparse—efficient in practice, e.g., NLP problems.
- Can use a hash function:  $\text{template} \rightarrow \{1, 2, \dots, d\}$ .

Ingredients in multiclass classification:

- Scoring functions for each class (similar to ranking).
- Represent labels in the input space  $\implies$  single weight vector.

We've seen

- How to generalize the perceptron algorithm to multiclass setting.
- Very simple idea. Was popular in NLP for structured prediction (e.g., tagging, parsing).

Next,

- How to generalize SVM to the multiclass setting.
- **Concept check:** Why might one prefer SVM / perceptron?



# Margin for Multiclass

Binary • Margin for  $(x^{(n)}, y^{(n)})$ :

$$y^{(n)} w^T x^{(n)} \quad (8)$$

- Want margin to be large and positive ( $w^T x^{(n)}$  has same sign as  $y^{(n)}$ )

Multiclass • Class-specific margin for  $(x^{(n)}, y^{(n)})$ :

$$h(x^{(n)}, y^{(n)}) - h(x^{(n)}, y). \quad (9)$$

- Difference between scores of the correct class and each other class
- Want margin to be large and positive for all  $y \neq y^{(n)}$ .

# Multiclass SVM: separable case

## Binary

$$\min_w \quad \frac{1}{2} \|w\|^2 \quad (10)$$

$$\text{s.t.} \quad \underbrace{y^{(n)} w^T x^{(n)}}_{\text{margin}} \geq 1 \quad \forall (x^{(n)}, y^{(n)}) \in \mathcal{D} \quad (11)$$

**Multiclass** As in the binary case, take 1 as our target margin.

$$m_{n,y}(w) \stackrel{\text{def}}{=} \underbrace{\langle w, \Psi(x^{(n)}, y^{(n)}) \rangle}_{\text{score of correct class}} - \underbrace{\langle w, \Psi(x^{(n)}, y) \rangle}_{\text{score of other class}} \quad (12)$$

$$\min_w \quad \frac{1}{2} \|w\|^2 \quad (13)$$

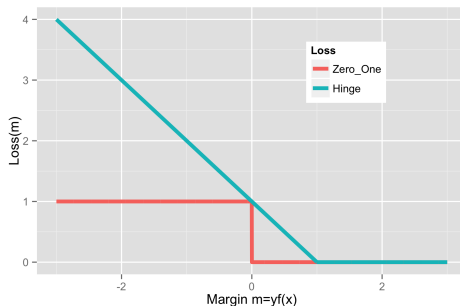
$$\text{s.t.} \quad m_{n,y}(w) \geq 1 \quad \forall (x^{(n)}, y^{(n)}) \in \mathcal{D}, y \neq y^{(n)} \quad (14)$$

**Exercise:** write the objective for the non-separable case

## Recap: hinge loss for binary classification

- Hinge loss: a convex upperbound on the 0-1 loss

$$\ell_{\text{hinge}}(y, \hat{y}) = \max(0, 1 - yh(x)) \quad (15)$$



# Generalized hinge loss

- What's the zero-one loss for multiclass classification?

$$\Delta(y, y') = \mathbb{I}\{y \neq y'\} \quad (16)$$

- In general, can also have different cost for each class.
- Upper bound on  $\Delta(y, y')$ .

$$\hat{y} \stackrel{\text{def}}{=} \arg \max_{y' \in \mathcal{Y}} \langle w, \Psi(x, y') \rangle \quad (17)$$

$$\implies \langle w, \Psi(x, y) \rangle \leq \langle w, \Psi(x, \hat{y}) \rangle \quad (18)$$

$$\implies \Delta(y, \hat{y}) \leq \Delta(y, \hat{y}) - \langle w, (\Psi(x, y) - \Psi(x, \hat{y})) \rangle \quad \text{When are they equal?} \quad (19)$$

- Generalized hinge loss:

$$\ell_{\text{hinge}}(y, x, w) \stackrel{\text{def}}{=} \max_{y' \in \mathcal{Y}} (\Delta(y, y') - \langle w, (\Psi(x, y) - \Psi(x, y')) \rangle) \quad (20)$$

# Multiclass SVM with Hinge Loss

- Recall the hinge loss formulation for binary SVM (without the bias term):

$$\min_{w \in \mathbb{R}^d} \frac{1}{2} \|w\|^2 + C \sum_{n=1}^N \max \left( 0, 1 - \underbrace{y^{(n)} w^T x^{(n)}}_{\text{margin}} \right).$$

- The multiclass objective:

$$\min_{w \in \mathbb{R}^d} \frac{1}{2} \|w\|^2 + C \sum_{n=1}^N \max_{y' \in \mathcal{Y}} \left( \Delta(y, y') - \underbrace{\langle w, (\Psi(x, y) - \Psi(x, y')) \rangle}_{\text{margin}} \right)$$

- $\Delta(y, y')$  as **target margin** for each class.
- If margin  $m_{n, y'}(w)$  meets or exceeds its target  $\Delta(y^{(n)}, y') \forall y \in \mathcal{Y}$ , then no loss on example  $n$ .

# Recap: What Have We Got?

- Problem: Multiclass classification  $\mathcal{Y} = \{1, \dots, k\}$
- Solution 1: One-vs-All
  - Train  $k$  models:  $h_1(x), \dots, h_k(x) : \mathcal{X} \rightarrow \mathbb{R}$ .
  - Predict with  $\arg \max_{y \in \mathcal{Y}} h_y(x)$ .
  - Gave simple example where this fails for linear classifiers
- Solution 2: Multiclass loss
  - Train one model:  $h(x, y) : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}$ .
  - Prediction involves solving  $\arg \max_{y \in \mathcal{Y}} h(x, y)$ .

# Does it work better in practice?

- Paper by Rifkin & Klautau: “In Defense of One-Vs-All Classification” (2004)
  - Extensive experiments, carefully done
    - albeit on relatively small UCI datasets
  - Suggests one-vs-all works just as well in practice
    - (or at least, the advantages claimed by earlier papers for multiclass methods were not compelling)
- Compared
  - many multiclass frameworks (including the one we discuss)
  - one-vs-all for SVMs with RBF kernel
  - one-vs-all for square loss with RBF kernel (for classification!)
- All performed roughly the same

# Why Are We Bothering with Multiclass?

- The framework we have developed for multiclass
  - compatibility features / scoring functions
  - multiclass margin
  - target margin / multiclass loss
- Generalizes to situations where  $k$  is very large and one-vs-all is intractable.
- Key idea is that we can generalize across outputs  $y$  by using features of  $y$ .



# Introduction to Structured Prediction

## Example: Part-of-speech (POS) Tagging

- Given a sentence, give a part of speech tag for each word:

$x$	$\underbrace{[\text{START}]}_{x_0}$	$\underbrace{\text{He}}_{x_1}$	$\underbrace{\text{eats}}_{x_2}$	$\underbrace{\text{apples}}_{x_3}$
$y$	$\underbrace{[\text{START}]}_{y_0}$	$\underbrace{\text{Pronoun}}_{y_1}$	$\underbrace{\text{Verb}}_{y_2}$	$\underbrace{\text{Noun}}_{y_3}$

- $\mathcal{V} = \{\text{all English words}\} \cup \{[\text{START}], ", ."]\}$
- $\mathcal{X} = \mathcal{V}^n, n = 1, 2, 3, \dots$  [Word sequences of any length]
- $\mathcal{P} = \{\text{START, Pronoun, Verb, Noun, Adjective}\}$
- $\mathcal{Y} = \mathcal{P}^n, n = 1, 2, 3, \dots$  [Part of speech sequence of any length]

# Multiclass Hypothesis Space

- **Discrete** output space:  $\mathcal{Y}(x)$ 
  - Very large but has structure, e.g., linear chain (sequence labeling), tree (parsing)
  - Size depends on input  $x$
- Base Hypothesis Space:  $\mathcal{H} = \{h : \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}\}$ 
  - $h(x, y)$  gives **compatibility score** between input  $x$  and output  $y$
- Multiclass hypothesis space

$$\mathcal{F} = \left\{ x \mapsto \arg \max_{y \in \mathcal{Y}} h(x, y) \mid h \in \mathcal{H} \right\}$$

- Final prediction function is an  $f \in \mathcal{F}$ .
- For each  $f \in \mathcal{F}$  there is an underlying compatibility score function  $h \in \mathcal{H}$ .

# Structured Prediction

- Part-of-speech tagging

$x$ :	he	eats	apples
$y$ :	pronoun	verb	noun

- Multiclass hypothesis space:

$$h(x, y) = w^T \Psi(x, y) \quad (21)$$

$$\mathcal{F} = \left\{ x \mapsto \arg \max_{y \in \mathcal{Y}} h(x, y) \mid h \in \mathcal{H} \right\} \quad (22)$$

- A special case of multiclass classification
- How to design the feature map  $\Psi$ ? What are the considerations?

# Unary features

- A **unary feature** only depends on
  - the label at a **single position**,  $y_i$ , and  $x$
- Example:

$$\phi_1(x, y_i) = \mathbb{1}[x_i = \text{runs}] \mathbb{1}[y_i = \text{Verb}]$$

$$\phi_2(x, y_i) = \mathbb{1}[x_i = \text{runs}] \mathbb{1}[y_i = \text{Noun}]$$

$$\phi_3(x, y_i) = \mathbb{1}[x_{i-1} = \text{He}] \mathbb{1}[x_i = \text{runs}] \mathbb{1}[y_i = \text{Verb}]$$

# Markov features

- A **markov feature** only depends on
  - two **adjacent** labels,  $y_{i-1}$  and  $y_i$ , and  $x$
- Example:

$$\theta_1(x, y_{i-1}, y_i) = \mathbb{1}[y_{i-1} = \text{Pronoun}] \mathbb{1}[y_i = \text{Verb}]$$

$$\theta_2(x, y_{i-1}, y_i) = \mathbb{1}[y_{i-1} = \text{Pronoun}] \mathbb{1}[y_i = \text{Noun}]$$

- Reminiscent of Markov models in the output space
- Possible to have higher-order features

## Local Feature Vector and Compatibility Score

- At each position  $i$  in sequence, define the **local feature vector** (unary and markov):

$$\Psi_i(x, y_{i-1}, y_i) = (\phi_1(x, y_i), \phi_2(x, y_i), \dots, \theta_1(x, y_{i-1}, y_i), \theta_2(x, y_{i-1}, y_i), \dots)$$

- And **local compatibility score** at position  $i$ :  $\langle w, \Psi_i(x, y_{i-1}, y_i) \rangle$ .
- The compatibility score for  $(x, y)$  is the sum of local compatibility scores:

$$\sum_i \langle w, \Psi_i(x, y_{i-1}, y_i) \rangle = \left\langle w, \sum_i \Psi_i(x, y_{i-1}, y_i) \right\rangle = \langle w, \Psi(x, y) \rangle, \quad (23)$$

where we define the **sequence feature vector** by

$$\Psi(x, y) = \sum_i \Psi_i(x, y_{i-1}, y_i). \quad \text{decomposable}$$

# Structured perceptron

Given a dataset  $\mathcal{D} = \{(x, y)\}$ ;

Initialize  $w \leftarrow 0$ ;

**for**  $iter = 1, 2, \dots, T$  **do**

**for**  $(x, y) \in \mathcal{D}$  **do**

$\hat{y} = \arg \max_{y' \in \mathcal{Y}(x)} w^T \psi(x, y')$ ;

**if**  $\hat{y} \neq y$  **then** // We've made a mistake

$w \leftarrow w + \Psi(x, y)$  ; // Move the scorer towards  $\psi(x, y)$

$w \leftarrow w - \Psi(x, \hat{y})$  ; // Move the scorer away from  $\psi(x, \hat{y})$

**end**

**end**

**end**

Identical to the multiclass perceptron algorithm except the  $\arg \max$  is now over the structured output space  $\mathcal{Y}(x)$ .



# Structured hinge loss

- Recall the generalized hinge loss

$$\ell_{\text{hinge}}(y, \hat{y}) \stackrel{\text{def}}{=} \max_{y' \in \mathcal{Y}(\mathbf{x})} (\Delta(y, y') + \langle w, (\Psi(\mathbf{x}, y') - \Psi(\mathbf{x}, y)) \rangle) \quad (24)$$

- What is  $\Delta(y, y')$  for two sequences?
- Hamming loss** is common:

$$\Delta(y, y') = \frac{1}{L} \sum_{i=1}^L \mathbb{1}[y_i \neq y'_i]$$

where  $L$  is the sequence length.

## Exercise:


- Write down the objective of structured SVM using the structured hinge loss.
- Stochastic sub-gradient descent for structured SVM (similar to HW3 P3)
- Compare with the structured perceptron algorithm

# The argmax problem for sequences

**Problem** To compute predictions, we need to find  $\arg\max_{y \in \mathcal{Y}(x)} \langle w, \Psi(x, y) \rangle$ , and  $|\mathcal{Y}(x)|$  is exponentially large.

**Observation**  $\Psi(x, y)$  decomposes to  $\sum_i \Psi_i(x, y)$ .

**Solution** Dynamic programming (similar to the Viterbi algorithm)




What's the running time?

# Conditional random field (CRF)



- Recall that we can write logistic regression in a general form:

$$p(y|x) = \frac{1}{Z(x)} \exp(w^\top \psi(x, y)).$$

- $Z$  is normalization constant:  $Z(x) = \sum_{y \in \mathcal{Y}} \exp(w^\top \psi(x, y))$ .
- Example: linear chain  $\{y_t\}$
- We can incorporate unary and Markov features:  $p(y|x) = \frac{1}{Z(x)} \exp(\sum_t w^\top \psi(x, y_t, y_{t-1}))$



Logistic Regression



Linear-chain CRFs

# Conditional random field (CRF)

- Compared to Structured SVM, CRF has a probabilistic interpretation.
- We can draw samples in the output space.
- How do we learn  $w$ ? Maximum log likelihood, and regularization term:  $\lambda \|w\|^2$
- Loss function:

$$\begin{aligned} l(w) &= -\frac{1}{N} \sum_{i=1}^N \log p(y^{(i)} | x^{(i)}) + \frac{1}{2} \lambda \|w\|^2 \\ &= -\frac{1}{N} \sum_i \sum_t \sum_k w_k \psi_k(y_t^{(i)}, y_{t-1}^{(i)}) + \frac{1}{N} \sum_i \log Z(x^{(i)}) + \frac{1}{2} \sum_k \lambda w_k^2 \end{aligned}$$

# Conditional random field (CRF)

- Loss function:

$$l(w) = -\frac{1}{N} \sum_i \sum_t \sum_k w_k \psi_k(x^{(i)}, y_t^{(i)}, y_{t-1}^{(i)}) + \frac{1}{N} \sum_i \log Z(x^{(i)}) + \frac{1}{2} \sum_k \lambda w_k^2$$

- Gradient:

$$\frac{\partial l(w)}{\partial w_k} = -\frac{1}{N} \sum_i \sum_t \sum_k \psi_k(x^{(i)}, y_t^{(i)}, y_{t-1}^{(i)}) \quad (25)$$

$$+ \frac{1}{N} \sum_i \frac{\partial}{\partial w_k} \log \sum_{y' \in Y} \exp\left(\sum_t \sum_{k'} w_{k'} \psi_{k'}(x^{(i)}, y'_t, y'_{t-1})\right) + \sum_k \lambda w_k \quad (26)$$

# Conditional random field (CRF)

- What is  $\frac{1}{N} \sum_i \sum_t \sum_k \psi_k(x^{(i)}, y_t^{(i)}, y_{t-1}^{(i)})$ ?
- It is the expectation  $\psi_k(x^{(i)}, y_t, y_{t-1})$  under the empirical distribution  $\tilde{p}(x, y) = \frac{1}{N} \sum_i \mathbb{1}[x = x^{(i)}] \mathbb{1}[y = y^{(i)}]$ .

## Conditional random field (CRF)

- What is  $\frac{1}{N} \sum_i \frac{\partial}{\partial w_k} \log \sum_{y' \in Y} \exp(\sum_t \sum_{k'} w_{k'} \psi_{k'}(x^{(i)}, y'_t, y'_{t-1}))$ ?

$$\frac{1}{N} \sum_i \frac{\partial}{\partial w_k} \log \sum_{y' \in Y} \exp(\sum_t \sum_{k'} w_{k'} \psi_{k'}(x^{(i)}, y'_t, y'_{t-1})) \quad (27)$$

$$= \frac{1}{N} \sum_i \left[ \sum_{y' \in Y} \exp(\sum_t \sum_{k'} w_{k'} \psi_{k'}(x^{(i)}, y'_t, y'_{t-1})) \right]^{-1} \quad (28)$$

$$\left[ \sum_{y' \in Y} \exp(\sum_t \sum_{k'} w_{k'} \psi_{k'}(x^{(i)}, y'_t, y'_{t-1})) \sum_t \psi_k(x^{(i)}, y'_t, y'_{t-1}) \right] \quad (29)$$

$$= \frac{1}{N} \sum_i \sum_t \sum_{y' \in Y} p(y'_t, y'_{t-1} | x) \psi_k(x^{(i)}, y'_t, y'_{t-1}) \quad (30)$$

- It is the expectation of  $\psi_k(x^{(i)}, y'_t, y'_{t-1})$  under the model distribution  $p(y'_t, y'_{t-1} | x)$ .



# Conditional random field (CRF)

- To compute the gradient, we need to infer expectation under the model distribution  $p(y|x)$ .
- Compare the learning algorithms: in structured SVM we need to compute the argmax, whereas in CRF we need to compute the model expectation.
- Both problems are NP-hard for general graphs.

- In the linear chain structure, we can use the forward-backward algorithm for inference, similar to Viterbi.
- Initiate  $\alpha_j(1) = \exp(w^\top \psi(y_1 = j, x_1))$
- Recursion:  $\alpha_j(t) = \sum_i \alpha_i(t-1) \exp(w^\top \psi(y_t = j, y_{t-1} = i, x_t))$
- Result:  $Z(x) = \sum_j \alpha_j(T)$
- Similar for the backward direction.
- Test time, again use Viterbi algorithm to infer argmax.
- The inference algorithm can be generalized to belief propagation (BP) in a tree structure (exact inference).
- In general graphs, we rely on approximate inference (e.g. loopy belief propagation).

- POS tag Relationship between constituents, e.g. NP is likely to be followed by a VP.
- Semantic segmentation  
Relationship between pixels, e.g. a grass pixel is likely to be next to another grass pixel, and a sky pixel is likely to be above a grass pixel.
- Multi-label learning  
An image may contain multiple class labels, e.g. a bus is likely to co-occur with a car.

## Multiclass algorithms

- Reduce to binary classification, e.g., OvA, AvA
  - Good enough for simple multiclass problems
  - They don't scale and have simplified assumptions
- Generalize binary classification algorithms using multiclass loss
  - Multi-class perceptron, multi-class logistics regression, multi-class SVM
- Structured prediction: Structured SVM, CRF. Data containing structure. Extremely large output space. Text and image applications.