

# Decision Trees and Boosting

Mengye Ren

NYU

Nov 14, 2023

# Review: Decision Trees

- Non-linear, non-metric, and non-parametric.
- Regression or classification.

# Review: Decision Trees

- Non-linear, non-metric, and non-parametric.
- Regression or classification.
- Interpretable, up to certain depth.

# Review: Decision Trees

- Non-linear, non-metric, and non-parametric.
- Regression or classification.
- Interpretable, up to certain depth.
- Greedy algorithm – maximizing the purity of nodes.

# Review: Decision Trees

- Non-linear, non-metric, and non-parametric.
- Regression or classification.
- Interpretable, up to certain depth.
- Greedy algorithm – maximizing the purity of nodes.
- Can overfit – need to limit the capacity.

# Bagging and Random Forests

## Recap: Statistics and Point Estimators

- We observe data  $\mathcal{D} = (x_1, x_2, \dots, x_n)$  sampled i.i.d. from a parametric distribution  $p(\cdot | \theta)$

## Recap: Statistics and Point Estimators

- We observe data  $\mathcal{D} = (x_1, x_2, \dots, x_n)$  sampled i.i.d. from a parametric distribution  $p(\cdot | \theta)$
- A **statistic**  $s = s(\mathcal{D})$  is any function of the data:



## Recap: Statistics and Point Estimators

- We observe data  $\mathcal{D} = (x_1, x_2, \dots, x_n)$  sampled i.i.d. from a parametric distribution  $p(\cdot | \theta)$
- A **statistic**  $s = s(\mathcal{D})$  is any function of the data:
  - E.g., sample mean, sample variance, histogram, empirical data distribution

## Recap: Statistics and Point Estimators

- We observe data  $\mathcal{D} = (x_1, x_2, \dots, x_n)$  sampled i.i.d. from a parametric distribution  $p(\cdot | \theta)$
- A **statistic**  $s = s(\mathcal{D})$  is any function of the data:
  - E.g., sample mean, sample variance, histogram, empirical data distribution
- A statistic  $\hat{\theta} = \hat{\theta}(\mathcal{D})$  is a **point estimator** of  $\theta$  if  $\hat{\theta} \approx \theta$

## Recap: Bias and Variance of an Estimator

- Statistics are random, so they have probability distributions.
- The distribution of a statistic is called a **sampling distribution**.
- The standard deviation of the sampling distribution is called the **standard error**.

## Recap: Bias and Variance of an Estimator

- Statistics are random, so they have probability distributions.
- The distribution of a statistic is called a **sampling distribution**.
- The standard deviation of the sampling distribution is called the **standard error**.
- Some parameters of the sampling distribution we might be interested in:

**Bias**  $\text{Bias}(\hat{\theta}) \stackrel{\text{def}}{=} \mathbb{E}[\hat{\theta}] - \theta.$

**Variance**  $\text{Var}(\hat{\theta}) \stackrel{\text{def}}{=} \mathbb{E}[\hat{\theta}^2] - \mathbb{E}^2[\hat{\theta}].$

## Recap: Bias and Variance of an Estimator

- Statistics are random, so they have probability distributions.
- The distribution of a statistic is called a **sampling distribution**.
- The standard deviation of the sampling distribution is called the **standard error**.
- Some parameters of the sampling distribution we might be interested in:

**Bias**  $\text{Bias}(\hat{\theta}) \stackrel{\text{def}}{=} \mathbb{E}[\hat{\theta}] - \theta.$

**Variance**  $\text{Var}(\hat{\theta}) \stackrel{\text{def}}{=} \mathbb{E}[\hat{\theta}^2] - \mathbb{E}^2[\hat{\theta}].$

- Why does variance matter if an estimator is unbiased?

## Recap: Bias and Variance of an Estimator

- Statistics are random, so they have probability distributions.
- The distribution of a statistic is called a **sampling distribution**.
- The standard deviation of the sampling distribution is called the **standard error**.
- Some parameters of the sampling distribution we might be interested in:

**Bias**  $\text{Bias}(\hat{\theta}) \stackrel{\text{def}}{=} \mathbb{E}[\hat{\theta}] - \theta.$

**Variance**  $\text{Var}(\hat{\theta}) \stackrel{\text{def}}{=} \mathbb{E}[\hat{\theta}^2] - \mathbb{E}^2[\hat{\theta}].$

- Why does variance matter if an estimator is unbiased?

## Recap: Bias and Variance of an Estimator

- Statistics are random, so they have probability distributions.
- The distribution of a statistic is called a **sampling distribution**.
- The standard deviation of the sampling distribution is called the **standard error**.
- Some parameters of the sampling distribution we might be interested in:

**Bias**  $\text{Bias}(\hat{\theta}) \stackrel{\text{def}}{=} \mathbb{E}[\hat{\theta}] - \theta.$

**Variance**  $\text{Var}(\hat{\theta}) \stackrel{\text{def}}{=} \mathbb{E}[\hat{\theta}^2] - \mathbb{E}^2[\hat{\theta}].$

- Why does variance matter if an estimator is unbiased?
  - $\hat{\theta}(\mathcal{D}) = x_1$  is an unbiased estimator of the mean of a Gaussian, but would be farther away from  $\theta$  than the sample mean.

## Variance of a Mean

- Let  $\hat{\theta}(\mathcal{D})$  be an unbiased estimator with variance  $\sigma^2$ :  $\mathbb{E}[\hat{\theta}] = \theta$ ,  $\text{Var}(\hat{\theta}) = \sigma^2$ .
- So far we have used a single statistic  $\hat{\theta} = \hat{\theta}(\mathcal{D})$  to estimate  $\theta$ .



## Variance of a Mean

- Let  $\hat{\theta}(\mathcal{D})$  be an unbiased estimator with variance  $\sigma^2$ :  $\mathbb{E}[\hat{\theta}] = \theta$ ,  $\text{Var}(\hat{\theta}) = \sigma^2$ .
- So far we have used a single statistic  $\hat{\theta} = \hat{\theta}(\mathcal{D})$  to estimate  $\theta$ .
- Its standard error is  $\sqrt{\text{Var}(\hat{\theta})} = \sigma$

## Variance of a Mean

- Let  $\hat{\theta}(\mathcal{D})$  be an unbiased estimator with variance  $\sigma^2$ :  $\mathbb{E}[\hat{\theta}] = \theta$ ,  $\text{Var}(\hat{\theta}) = \sigma^2$ .
- So far we have used a single statistic  $\hat{\theta} = \hat{\theta}(\mathcal{D})$  to estimate  $\theta$ .
- Its standard error is  $\sqrt{\text{Var}(\hat{\theta})} = \sigma$
- Consider a new estimator that takes the average of i.i.d.  $\hat{\theta}_1, \dots, \hat{\theta}_n$  where  $\hat{\theta}_i = \hat{\theta}(\mathcal{D}^i)$ .

## Variance of a Mean

- Let  $\hat{\theta}(\mathcal{D})$  be an unbiased estimator with variance  $\sigma^2$ :  $\mathbb{E}[\hat{\theta}] = \theta$ ,  $\text{Var}(\hat{\theta}) = \sigma^2$ .
- So far we have used a single statistic  $\hat{\theta} = \hat{\theta}(\mathcal{D})$  to estimate  $\theta$ .
- Its standard error is  $\sqrt{\text{Var}(\hat{\theta})} = \sigma$
- Consider a new estimator that takes the average of i.i.d.  $\hat{\theta}_1, \dots, \hat{\theta}_n$  where  $\hat{\theta}_i = \hat{\theta}(\mathcal{D}^i)$ .

## Variance of a Mean

- Let  $\hat{\theta}(\mathcal{D})$  be an unbiased estimator with variance  $\sigma^2$ :  $\mathbb{E}[\hat{\theta}] = \theta$ ,  $\text{Var}(\hat{\theta}) = \sigma^2$ .
- So far we have used a single statistic  $\hat{\theta} = \hat{\theta}(\mathcal{D})$  to estimate  $\theta$ .
- Its standard error is  $\sqrt{\text{Var}(\hat{\theta})} = \sigma$
- Consider a new estimator that takes the average of i.i.d.  $\hat{\theta}_1, \dots, \hat{\theta}_n$  where  $\hat{\theta}_i = \hat{\theta}(\mathcal{D}^i)$ .
- The average has the same expected value but smaller standard error (recall that  $\text{Var}(cX) = c^2 \text{Var}(X)$ , and that the  $\hat{\theta}_i$ -s are uncorrelated):

$$\mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n \hat{\theta}_i\right] = \theta \quad \text{Var}\left[\frac{1}{n} \sum_{i=1}^n \hat{\theta}_i\right] = \frac{\sigma^2}{n} \quad (1)$$

# Averaging Independent Prediction Functions

- Suppose we have  $B$  independent training sets, all drawn from the same distribution ( $\mathcal{D} \sim p(\cdot | \theta)$ ).

# Averaging Independent Prediction Functions

- Suppose we have  $B$  independent training sets, all drawn from the same distribution ( $\mathcal{D} \sim p(\cdot | \theta)$ ).
- Our learning algorithm gives us  $B$  prediction functions:  $\hat{f}_1(x), \hat{f}_2(x), \dots, \hat{f}_B(x)$

# Averaging Independent Prediction Functions

- Suppose we have  $B$  independent training sets, all drawn from the same distribution ( $\mathcal{D} \sim p(\cdot \mid \theta)$ ).
- Our learning algorithm gives us  $B$  prediction functions:  $\hat{f}_1(x), \hat{f}_2(x), \dots, \hat{f}_B(x)$
- We will define the average prediction function as:

$$\hat{f}_{\text{avg}} \stackrel{\text{def}}{=} \frac{1}{B} \sum_{b=1}^B \hat{f}_b \quad (2)$$

# Averaging Reduces Variance of Predictions

- The average prediction for  $x_0$  is

$$\hat{f}_{\text{avg}}(x_0) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b(x_0).$$

- $\hat{f}_{\text{avg}}(x_0)$  and  $\hat{f}_b(x_0)$  have the same expected value, but
- $\hat{f}_{\text{avg}}(x_0)$  has smaller variance:

$$\text{Var}(\hat{f}_{\text{avg}}(x_0)) = \frac{1}{B} \text{Var}(\hat{f}_1(x_0))$$



# Averaging Reduces Variance of Predictions

- The average prediction for  $x_0$  is

$$\hat{f}_{\text{avg}}(x_0) = \frac{1}{B} \sum_{b=1}^B \hat{f}_b(x_0).$$

- $\hat{f}_{\text{avg}}(x_0)$  and  $\hat{f}_b(x_0)$  have the same expected value, but
- $\hat{f}_{\text{avg}}(x_0)$  has smaller variance:

$$\text{Var}(\hat{f}_{\text{avg}}(x_0)) = \frac{1}{B} \text{Var}(\hat{f}_1(x_0))$$

- **Problem:** in practice we don't have  $B$  independent training sets!

# The Bootstrap Sample

How do we simulate multiple samples when we only have one?

- A **bootstrap sample** from  $\mathcal{D}_n = (x_1, \dots, x_n)$  is a sample of size  $n$  drawn *with replacement* from  $\mathcal{D}_n$

# The Bootstrap Sample

How do we simulate multiple samples when we only have one?

- A **bootstrap sample** from  $\mathcal{D}_n = (x_1, \dots, x_n)$  is a sample of size  $n$  drawn *with replacement* from  $\mathcal{D}_n$
- Some elements of  $\mathcal{D}_n$  will show up multiple times, and some won't show up at all

# The Bootstrap Sample

How do we simulate multiple samples when we only have one?

- A **bootstrap sample** from  $\mathcal{D}_n = (x_1, \dots, x_n)$  is a sample of size  $n$  drawn *with replacement* from  $\mathcal{D}_n$
- Some elements of  $\mathcal{D}_n$  will show up multiple times, and some won't show up at all
- Each  $x_i$  has a probability of  $(1 - 1/n)^n$  of not being included in a given bootstrap sample

# The Bootstrap Sample

How do we simulate multiple samples when we only have one?

- A **bootstrap sample** from  $\mathcal{D}_n = (x_1, \dots, x_n)$  is a sample of size  $n$  drawn *with replacement* from  $\mathcal{D}_n$
- Some elements of  $\mathcal{D}_n$  will show up multiple times, and some won't show up at all
- Each  $x_i$  has a probability of  $(1 - 1/n)^n$  of not being included in a given bootstrap sample
- For large  $n$ ,

$$\left(1 - \frac{1}{n}\right)^n \approx \frac{1}{e} \approx .368. \quad (3)$$

- So we expect ~63.2% of elements of  $\mathcal{D}_n$  will show up at least once.

# The Bootstrap Method

## Definition

A **bootstrap method** simulates  $B$  independent samples from  $P$  by taking  $B$  bootstrap samples from the sample  $\mathcal{D}_n$ .

# The Bootstrap Method

## Definition

A **bootstrap method** simulates  $B$  independent samples from  $P$  by taking  $B$  bootstrap samples from the sample  $\mathcal{D}_n$ .

- Given original data  $\mathcal{D}_n$ , compute  $B$  bootstrap samples  $D_n^1, \dots, D_n^B$ .

# The Bootstrap Method

## Definition

A **bootstrap method** simulates  $B$  independent samples from  $P$  by taking  $B$  bootstrap samples from the sample  $\mathcal{D}_n$ .

- Given original data  $\mathcal{D}_n$ , compute  $B$  bootstrap samples  $D_n^1, \dots, D_n^B$ .
- For each bootstrap sample, compute some function

$$\phi(D_n^1), \dots, \phi(D_n^B)$$



# The Bootstrap Method

## Definition

A **bootstrap method** simulates  $B$  independent samples from  $P$  by taking  $B$  bootstrap samples from the sample  $\mathcal{D}_n$ .

- Given original data  $\mathcal{D}_n$ , compute  $B$  bootstrap samples  $D_n^1, \dots, D_n^B$ .
- For each bootstrap sample, compute some function

$$\phi(D_n^1), \dots, \phi(D_n^B)$$

- Use these values as though  $D_n^1, \dots, D_n^B$  were i.i.d. samples from  $P$ .
- This often ends up being very close to what we'd get with independent samples from  $P$ !

# Independent Samples vs. Bootstrap Samples

- Point estimator  $\hat{\alpha} = \hat{\alpha}(\mathcal{D}_{100})$  for samples of size 100, for a synthetic case where the data generating distribution is known
- Histograms of  $\hat{\alpha}$  based on
  - 1000 independent samples of size 100 (left), vs.
  - 1000 bootstrap samples of size 100 (right)

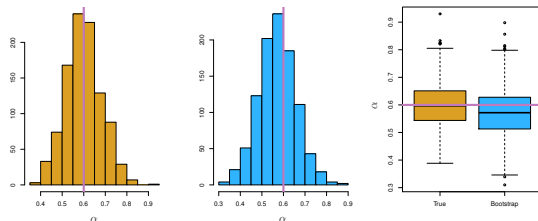


Figure 5.10 from *ISLR* (Springer, 2013) with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani.

## Key ideas:

- In general, **ensemble methods** combine multiple weak models into a single, more powerful model

## Key ideas:

- In general, **ensemble methods** combine multiple weak models into a single, more powerful model
- Averaging i.i.d. estimates reduces variance without changing bias

## Key ideas:

- In general, **ensemble methods** combine multiple weak models into a single, more powerful model
- Averaging i.i.d. estimates reduces variance without changing bias
- We can use bootstrap to simulate multiple data samples and average them

## Key ideas:

- In general, **ensemble methods** combine multiple weak models into a single, more powerful model
- Averaging i.i.d. estimates reduces variance without changing bias
- We can use bootstrap to simulate multiple data samples and average them
- Parallel ensemble (e.g., bagging): models are built independently

## Key ideas:

- In general, **ensemble methods** combine multiple weak models into a single, more powerful model
- Averaging i.i.d. estimates reduces variance without changing bias
- We can use bootstrap to simulate multiple data samples and average them
- Parallel ensemble (e.g., bagging): models are built independently
- Sequential ensemble (e.g., boosting): models are built sequentially

## Key ideas:

- In general, **ensemble methods** combine multiple weak models into a single, more powerful model
- Averaging i.i.d. estimates reduces variance without changing bias
- We can use bootstrap to simulate multiple data samples and average them
- Parallel ensemble (e.g., bagging): models are built independently
- Sequential ensemble (e.g., boosting): models are built sequentially
  - We try to find new learners that do well where previous learners fall short



# Bagging: Bootstrap Aggregation

- We draw  $B$  bootstrap samples  $D^1, \dots, D^B$  from original data  $\mathcal{D}$

# Bagging: Bootstrap Aggregation

- We draw  $B$  bootstrap samples  $D^1, \dots, D^B$  from original data  $\mathcal{D}$
- Let  $\hat{f}_1, \hat{f}_2, \dots, \hat{f}_B$  be the prediction functions resulting from training on  $D^1, \dots, D^B$ , respectively

# Bagging: Bootstrap Aggregation

- We draw  $B$  bootstrap samples  $D^1, \dots, D^B$  from original data  $\mathcal{D}$
- Let  $\hat{f}_1, \hat{f}_2, \dots, \hat{f}_B$  be the prediction functions resulting from training on  $D^1, \dots, D^B$ , respectively
- The **bagged prediction function** is a *combination* of these:

$$\hat{f}_{\text{avg}}(x) = \text{Combine} \left( \hat{f}_1(x), \hat{f}_2(x), \dots, \hat{f}_B(x) \right)$$

# Bagging: Bootstrap Aggregation

- Bagging is a general method for variance reduction, but it is particularly useful for decision trees

# Bagging: Bootstrap Aggregation

- Bagging is a general method for variance reduction, but it is particularly useful for decision trees
- For classification, averaging doesn't make sense; we can take a **majority vote** instead

# Bagging: Bootstrap Aggregation

- Bagging is a general method for variance reduction, but it is particularly useful for decision trees
- For classification, averaging doesn't make sense; we can take a **majority vote** instead
- Increasing the number of trees we use in bagging does not lead to overfitting

# Bagging: Bootstrap Aggregation

- Bagging is a general method for variance reduction, but it is particularly useful for decision trees
- For classification, averaging doesn't make sense; we can take a **majority vote** instead
- Increasing the number of trees we use in bagging does not lead to overfitting
- Is there a downside, compared to having a single decision tree?

# Bagging: Bootstrap Aggregation

- Bagging is a general method for variance reduction, but it is particularly useful for decision trees
- For classification, averaging doesn't make sense; we can take a **majority vote** instead
- Increasing the number of trees we use in bagging does not lead to overfitting
- Is there a downside, compared to having a single decision tree?
- Yes: if we have many trees, the bagged predictor is much less interpretable



## Aside: Out-of-Bag Error Estimation

- Recall that each bagged predictor was trained on about 63% of the data.
- The remaining 37% are called **out-of-bag (OOB)** observations.

## Aside: Out-of-Bag Error Estimation

- Recall that each bagged predictor was trained on about 63% of the data.
- The remaining 37% are called **out-of-bag (OOB)** observations.
- For  $i$ th training point, let

$$S_i = \{b \mid D^b \text{ does not contain } i\text{th point}\}$$

- The **OOB prediction** on  $x_i$  is

$$\hat{f}_{\text{OOB}}(x_i) = \frac{1}{|S_i|} \sum_{b \in S_i} \hat{f}_b(x_i)$$

## Aside: Out-of-Bag Error Estimation

- Recall that each bagged predictor was trained on about 63% of the data.
- The remaining 37% are called **out-of-bag (OOB)** observations.
- For  $i$ th training point, let

$$S_i = \{b \mid D^b \text{ does not contain } i\text{th point}\}$$

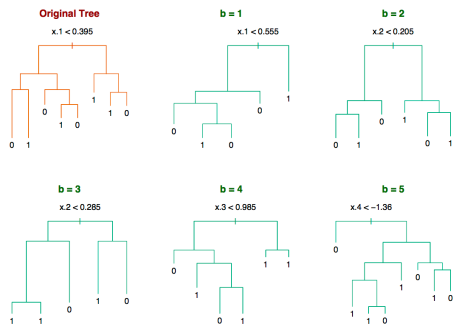
- The **OOB prediction** on  $x_i$  is

$$\hat{f}_{\text{OOB}}(x_i) = \frac{1}{|S_i|} \sum_{b \in S_i} \hat{f}_b(x_i)$$

- The OOB error is a good estimate of the test error
- Similar to cross validation error: both are computed on the training set

# Applying Bagging to Classification Trees


- Input space  $\mathcal{X} = \mathbb{R}^5$  and output space  $\mathcal{Y} = \{-1, 1\}$ . Sample size  $n = 30$ .



From HTF Figure 8.9

# Applying Bagging to Classification Trees

- Input space  $\mathcal{X} = \mathbb{R}^5$  and output space  $\mathcal{Y} = \{-1, 1\}$ . Sample size  $n = 30$ .




- Each bootstrap tree is quite different: different splitting variable at the root!

From HTF Figure 8.9

# Applying Bagging to Classification Trees

- Input space  $\mathcal{X} = \mathbb{R}^5$  and output space  $\mathcal{Y} = \{-1, 1\}$ . Sample size  $n = 30$ .




- Each bootstrap tree is quite different: different splitting variable at the root!
- High variance:** small perturbations of the training data lead to a high degree of model variability

From HTF Figure 8.9

# Applying Bagging to Classification Trees

- Input space  $\mathcal{X} = \mathbb{R}^5$  and output space  $\mathcal{Y} = \{-1, 1\}$ . Sample size  $n = 30$ .



- Each bootstrap tree is quite different: different splitting variable at the root!
- High variance:** small perturbations of the training data lead to a high degree of model variability
- Bagging helps most when the base learners are relatively unbiased but have high variance (exactly the case for decision trees)

# Motivating Random Forests: Correlated Prediction Functions

Recall the motivating principle of bagging:

- For  $\hat{\theta}_1, \dots, \hat{\theta}_n$  *i.i.d.* with  $\mathbb{E}[\hat{\theta}] = \theta$  and  $\text{Var}[\hat{\theta}] = \sigma^2$ ,

$$\mathbb{E}\left[\frac{1}{n}\sum_{i=1}^n \hat{\theta}_i\right] = \mu \quad \text{Var}\left[\frac{1}{n}\sum_{i=1}^n \hat{\theta}_i\right] = \frac{\sigma^2}{n}.$$



# Motivating Random Forests: Correlated Prediction Functions

Recall the motivating principle of bagging:

- For  $\hat{\theta}_1, \dots, \hat{\theta}_n$  *i.i.d.* with  $\mathbb{E}[\hat{\theta}] = \theta$  and  $\text{Var}[\hat{\theta}] = \sigma^2$ ,

$$\mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n \hat{\theta}_i\right] = \mu \quad \text{Var}\left[\frac{1}{n} \sum_{i=1}^n \hat{\theta}_i\right] = \frac{\sigma^2}{n}.$$

- What if  $\hat{\theta}$ 's are correlated?

# Motivating Random Forests: Correlated Prediction Functions

Recall the motivating principle of bagging:

- For  $\hat{\theta}_1, \dots, \hat{\theta}_n$  *i.i.d.* with  $\mathbb{E}[\hat{\theta}] = \theta$  and  $\text{Var}[\hat{\theta}] = \sigma^2$ ,

$$\mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n \hat{\theta}_i\right] = \mu \quad \text{Var}\left[\frac{1}{n} \sum_{i=1}^n \hat{\theta}_i\right] = \frac{\sigma^2}{n}.$$

- What if  $\hat{\theta}$ 's are correlated?
- For large  $n$ , the covariance term dominates, limiting the benefits of averaging

# Motivating Random Forests: Correlated Prediction Functions

Recall the motivating principle of bagging:

- For  $\hat{\theta}_1, \dots, \hat{\theta}_n$  *i.i.d.* with  $\mathbb{E}[\hat{\theta}] = \theta$  and  $\text{Var}[\hat{\theta}] = \sigma^2$ ,

$$\mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n \hat{\theta}_i\right] = \mu \quad \text{Var}\left[\frac{1}{n} \sum_{i=1}^n \hat{\theta}_i\right] = \frac{\sigma^2}{n}.$$

- What if  $\hat{\theta}$ 's are correlated?
- For large  $n$ , the covariance term dominates, limiting the benefits of averaging
- Bootstrap samples are

# Motivating Random Forests: Correlated Prediction Functions

Recall the motivating principle of bagging:

- For  $\hat{\theta}_1, \dots, \hat{\theta}_n$  *i.i.d.* with  $\mathbb{E}[\hat{\theta}] = \theta$  and  $\text{Var}[\hat{\theta}] = \sigma^2$ ,

$$\mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n \hat{\theta}_i\right] = \mu \quad \text{Var}\left[\frac{1}{n} \sum_{i=1}^n \hat{\theta}_i\right] = \frac{\sigma^2}{n}.$$

- What if  $\hat{\theta}$ 's are correlated?
- For large  $n$ , the covariance term dominates, limiting the benefits of averaging
- Bootstrap samples are
  - independent samples from the training set, but

# Motivating Random Forests: Correlated Prediction Functions

Recall the motivating principle of bagging:

- For  $\hat{\theta}_1, \dots, \hat{\theta}_n$  *i.i.d.* with  $\mathbb{E}[\hat{\theta}] = \theta$  and  $\text{Var}[\hat{\theta}] = \sigma^2$ ,

$$\mathbb{E}\left[\frac{1}{n}\sum_{i=1}^n \hat{\theta}_i\right] = \mu \quad \text{Var}\left[\frac{1}{n}\sum_{i=1}^n \hat{\theta}_i\right] = \frac{\sigma^2}{n}.$$

- What if  $\hat{\theta}$ 's are correlated?
- For large  $n$ , the covariance term dominates, limiting the benefits of averaging
- Bootstrap samples are
  - independent samples from the training set, but
  - **not** independent samples from  $P_{\mathcal{X} \times \mathcal{Y}}$

# Motivating Random Forests: Correlated Prediction Functions

Recall the motivating principle of bagging:

- For  $\hat{\theta}_1, \dots, \hat{\theta}_n$  *i.i.d.* with  $\mathbb{E}[\hat{\theta}] = \theta$  and  $\text{Var}[\hat{\theta}] = \sigma^2$ ,

$$\mathbb{E}\left[\frac{1}{n}\sum_{i=1}^n \hat{\theta}_i\right] = \mu \quad \text{Var}\left[\frac{1}{n}\sum_{i=1}^n \hat{\theta}_i\right] = \frac{\sigma^2}{n}.$$

- What if  $\hat{\theta}$ 's are correlated?
- For large  $n$ , the covariance term dominates, limiting the benefits of averaging
- Bootstrap samples are
  - independent samples from the training set, but
  - **not** independent samples from  $P_{\mathcal{X} \times \mathcal{Y}}$
- Can we reduce the dependence between  $\hat{f}_i$ 's?

# Random Forests

## Key idea

Use bagged decision trees, but modify the tree-growing procedure to reduce the dependence between trees.

- Build a collection of trees independently (in parallel), as before

## Key idea

Use bagged decision trees, but modify the tree-growing procedure to reduce the dependence between trees.

- Build a collection of trees independently (in parallel), as before
- When constructing each tree node, restrict choice of splitting variable to a randomly chosen subset of features of size  $m$ 
  - This prevents a situation where all trees are dominated by the same small number of strong features (and are therefore too similar to each other)



## Key idea

Use bagged decision trees, but modify the tree-growing procedure to reduce the dependence between trees.


- Build a collection of trees independently (in parallel), as before
- When constructing each tree node, restrict choice of splitting variable to a randomly chosen subset of features of size  $m$ 
  - This prevents a situation where all trees are dominated by the same small number of strong features (and are therefore too similar to each other)
- We typically choose  $m \approx \sqrt{p}$ , where  $p$  is the number of features (or we can choose  $m$  using cross validation)

## Key idea

Use bagged decision trees, but modify the tree-growing procedure to reduce the dependence between trees.

- Build a collection of trees independently (in parallel), as before
- When constructing each tree node, restrict choice of splitting variable to a randomly chosen subset of features of size  $m$ 
  - This prevents a situation where all trees are dominated by the same small number of strong features (and are therefore too similar to each other)
- We typically choose  $m \approx \sqrt{p}$ , where  $p$  is the number of features (or we can choose  $m$  using cross validation)
- If  $m = p$ , this is just bagging

# Random Forests: Effect of $m$



From *An Introduction to Statistical Learning, with applications in R* (Springer, 2013) with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani.

- The usual approach is to build very deep trees—low bias but **high variance**

- The usual approach is to build very deep trees—low bias but **high variance**
- Ensembling many models reduces variance
  - Motivation: Mean of i.i.d. estimates has smaller variance than single estimate

- The usual approach is to build very deep trees—low bias but **high variance**
- Ensembling many models reduces variance
  - Motivation: Mean of i.i.d. estimates has smaller variance than single estimate
- Use bootstrap to simulate many data samples from one dataset
  - $\implies$  Bagged decision trees

- The usual approach is to build very deep trees—low bias but **high variance**
- Ensembling many models reduces variance
  - Motivation: Mean of i.i.d. estimates has smaller variance than single estimate
- Use bootstrap to simulate many data samples from one dataset
  - $\implies$  Bagged decision trees
- But bootstrap samples (and the induced models) are correlated

- The usual approach is to build very deep trees—low bias but **high variance**
- Ensembling many models reduces variance
  - Motivation: Mean of i.i.d. estimates has smaller variance than single estimate
- Use bootstrap to simulate many data samples from one dataset
  - $\implies$  Bagged decision trees
- But bootstrap samples (and the induced models) are correlated
- Ensembling works better when we combine a diverse set of prediction functions
  - $\implies$  Random forests: select a random subset of features for each decision tree



# Boosting

**Bagging** Reduce variance of a low bias, high variance estimator by ensembling many estimators trained in parallel (on different datasets obtained through sampling).

# Boosting: Overview

- Bagging** Reduce variance of a low bias, high variance estimator by ensembling many estimators trained in parallel (on different datasets obtained through sampling).
- Boosting** Reduce the error rate of a high bias estimator by ensembling many estimators trained in sequence (without bootstrapping).

# Boosting: Overview

**Bagging** Reduce variance of a low bias, high variance estimator by ensembling many estimators trained in parallel (on different datasets obtained through sampling).

**Boosting** Reduce the error rate of a high bias estimator by ensembling many estimators trained in sequence (without bootstrapping).

- Like bagging, boosting is a general method that is particularly popular with decision trees.

# Boosting: Overview

**Bagging** Reduce variance of a low bias, high variance estimator by ensembling many estimators trained in parallel (on different datasets obtained through sampling).

**Boosting** Reduce the error rate of a high bias estimator by ensembling many estimators trained in sequence (without bootstrapping).

- Like bagging, boosting is a general method that is particularly popular with decision trees.
- Main intuition: instead of fitting the data very closely using a large decision tree, train gradually, using a sequence of simpler trees

# Boosting: Overview

- A **weak/base learner** is a classifier that does slightly better than chance.
- Weak learners are like rules of thumb:
  - “Inheritance”  $\implies$  spam
  - From a friend  $\implies$  not spam

# Boosting: Overview

- A **weak/base learner** is a classifier that does slightly better than chance.
- Weak learners are like rules of thumb:
  - “Inheritance”  $\implies$  spam
  - From a friend  $\implies$  not spam
- Key idea:

# Boosting: Overview

- A **weak/base learner** is a classifier that does slightly better than chance.
- Weak learners are like rules of thumb:
  - “Inheritance”  $\implies$  spam
  - From a friend  $\implies$  not spam
- **Key idea:**
  - Each weak learner focuses on different training examples (*reweighted data*)



# Boosting: Overview

- A **weak/base learner** is a classifier that does slightly better than chance.
- Weak learners are like rules of thumb:
  - “Inheritance”  $\implies$  spam
  - From a friend  $\implies$  not spam
- **Key idea:**
  - Each weak learner focuses on different training examples (*reweighted data*)
  - Weak learners make different contributions to the final prediction (*reweighted classifier*)

# Boosting: Overview

- A **weak/base learner** is a classifier that does slightly better than chance.
- Weak learners are like rules of thumb:
  - “Inheritance”  $\implies$  spam
  - From a friend  $\implies$  not spam
- **Key idea:**
  - Each weak learner focuses on different training examples (*reweighted data*)
  - Weak learners make different contributions to the final prediction (*reweighted classifier*)
- A set of smaller, simpler trees may improve interpretability

# Boosting: Overview

- A **weak/base learner** is a classifier that does slightly better than chance.
- Weak learners are like rules of thumb:
  - “Inheritance”  $\implies$  spam
  - From a friend  $\implies$  not spam
- **Key idea:**
  - Each weak learner focuses on different training examples (*reweighted data*)
  - Weak learners make different contributions to the final prediction (*reweighted classifier*)
- A set of smaller, simpler trees may improve interpretability
- We'll focus on a specific implementation, AdaBoost (Freund & Schapire, 1997)

# AdaBoost: Setting

- Binary classification:  $\mathcal{Y} = \{-1, 1\}$

# AdaBoost: Setting

- Binary classification:  $\mathcal{Y} = \{-1, 1\}$
- Base hypothesis space  $\mathcal{H} = \{h : \mathcal{X} \rightarrow \{-1, 1\}\}$ .

# AdaBoost: Setting

- Binary classification:  $\mathcal{Y} = \{-1, 1\}$
- Base hypothesis space  $\mathcal{H} = \{h : \mathcal{X} \rightarrow \{-1, 1\}\}$ .
- Typical base hypothesis spaces:
  - **Decision stumps** (tree with a single split)
  - Trees with few terminal nodes
  - Linear decision functions

# Weighted Training Set

Each base learner is trained on weighted data.

- Training set  $\mathcal{D} = ((x_1, y_1), \dots, (x_n, y_n))$ .
- Weights  $(w_1, \dots, w_n)$  associated with each example.

# Weighted Training Set

Each base learner is trained on weighted data.

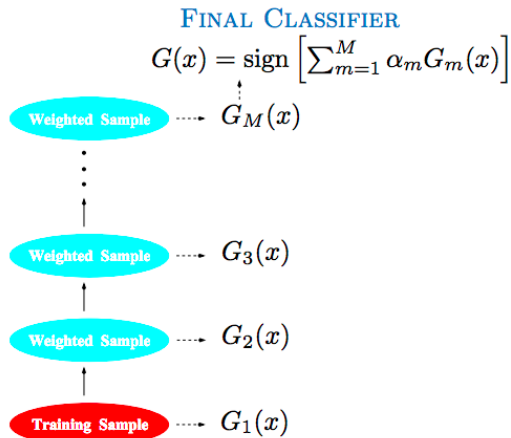
- Training set  $\mathcal{D} = ((x_1, y_1), \dots, (x_n, y_n))$ .
- Weights  $(w_1, \dots, w_n)$  associated with each example.
- **Weighted empirical risk:**

$$\hat{R}_n^w(f) \stackrel{\text{def}}{=} \frac{1}{W} \sum_{i=1}^n w_i \ell(f(x_i), y_i) \quad \text{where } W = \sum_{i=1}^n w_i$$

- Examples with larger weights affect the loss more.



# AdaBoost: Schematic



From ESL Figure 10.1

# AdaBoost: Sketch of the Algorithm

- Start with equal weights for all training points:  $w_1 = \dots = w_n = 1$

# AdaBoost: Sketch of the Algorithm

- Start with equal weights for all training points:  $w_1 = \dots = w_n = 1$
- Repeat for  $m = 1, \dots, M$  (where  $M$  is the number of classifiers we plan to train):

# AdaBoost: Sketch of the Algorithm

- Start with equal weights for all training points:  $w_1 = \dots = w_n = 1$
- Repeat for  $m = 1, \dots, M$  (where  $M$  is the number of classifiers we plan to train):
  - Train base classifier  $G_m(x)$  on the weighted training data; this classifier may not fit the data well

# AdaBoost: Sketch of the Algorithm

- Start with equal weights for all training points:  $w_1 = \dots = w_n = 1$
- Repeat for  $m = 1, \dots, M$  (where  $M$  is the number of classifiers we plan to train):
  - Train base classifier  $G_m(x)$  on the weighted training data; this classifier may not fit the data well
  - Increase the weight of the points misclassified by  $G_m(x)$  (this is the key idea of boosting!)

# AdaBoost: Sketch of the Algorithm

- Start with equal weights for all training points:  $w_1 = \dots = w_n = 1$
- Repeat for  $m = 1, \dots, M$  (where  $M$  is the number of classifiers we plan to train):
  - Train base classifier  $G_m(x)$  on the weighted training data; this classifier may not fit the data well
  - Increase the weight of the points misclassified by  $G_m(x)$  (this is the key idea of boosting!)
- Our final prediction is  $G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$

# AdaBoost: Classifier Weights

- Our final prediction is  $G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$ .
- We would like  $\alpha_m$  to be:
  - Nonnegative
  - Larger when  $G_m$  fits its weighted training data well
- The **weighted 0-1 error** of  $G_m(x)$  is

$$\text{err}_m = \frac{1}{W} \sum_{i=1}^n w_i \mathbb{1}[y_i \neq G_m(x_i)] \quad \text{where } W = \sum_{i=1}^n w_i.$$

- $\text{err}_m \in [0, 1]$

# AdaBoost: Classifier Weights

- The weight of classifier  $G_m(x)$  is  $\alpha_m = \ln \left( \frac{1 - \text{err}_m}{\text{err}_m} \right)$



# AdaBoost: Classifier Weights

- The weight of classifier  $G_m(x)$  is  $\alpha_m = \ln \left( \frac{1 - \text{err}_m}{\text{err}_m} \right)$



- Higher weighted error  $\implies$  lower weight

## AdaBoost: Example Reweighting

- We train  $G_m$  to minimize weighted error; the resulting error rate is  $\text{err}_m$
- Then  $\alpha_m = \ln\left(\frac{1-\text{err}_m}{\text{err}_m}\right)$  is the weight of  $G_m$  in the final ensemble

We want the next base learner to focus more on examples misclassified by the previous learner.

## AdaBoost: Example Reweighting

- We train  $G_m$  to minimize weighted error; the resulting error rate is  $\text{err}_m$
- Then  $\alpha_m = \ln\left(\frac{1-\text{err}_m}{\text{err}_m}\right)$  is the weight of  $G_m$  in the final ensemble

We want the next base learner to focus more on examples misclassified by the previous learner.

- Suppose  $w_i$  is the weight of example  $x_i$  before training:

## AdaBoost: Example Reweighting

- We train  $G_m$  to minimize weighted error; the resulting error rate is  $\text{err}_m$
- Then  $\alpha_m = \ln\left(\frac{1-\text{err}_m}{\text{err}_m}\right)$  is the weight of  $G_m$  in the final ensemble

We want the next base learner to focus more on examples misclassified by the previous learner.

- Suppose  $w_i$  is the weight of example  $x_i$  before training:
  - If  $G_m$  classifies  $x_i$  correctly, keep  $w_i$  as is

## AdaBoost: Example Reweighting

- We train  $G_m$  to minimize weighted error; the resulting error rate is  $\text{err}_m$
- Then  $\alpha_m = \ln\left(\frac{1-\text{err}_m}{\text{err}_m}\right)$  is the weight of  $G_m$  in the final ensemble

We want the next base learner to focus more on examples misclassified by the previous learner.

- Suppose  $w_i$  is the weight of example  $x_i$  before training:
  - If  $G_m$  classifies  $x_i$  correctly, keep  $w_i$  as is
  - Otherwise, increase  $w_i$ :

$$\begin{aligned}w_i &\leftarrow w_i e^{\alpha_m} \\ &= w_i \left( \frac{1 - \text{err}_m}{\text{err}_m} \right)\end{aligned}$$

## AdaBoost: Example Reweighting

- We train  $G_m$  to minimize weighted error; the resulting error rate is  $\text{err}_m$
- Then  $\alpha_m = \ln\left(\frac{1-\text{err}_m}{\text{err}_m}\right)$  is the weight of  $G_m$  in the final ensemble

We want the next base learner to focus more on examples misclassified by the previous learner.

- Suppose  $w_i$  is the weight of example  $x_i$  before training:
  - If  $G_m$  classifies  $x_i$  correctly, keep  $w_i$  as is
  - Otherwise, increase  $w_i$ :

$$\begin{aligned}w_i &\leftarrow w_i e^{\alpha_m} \\ &= w_i \left(\frac{1-\text{err}_m}{\text{err}_m}\right)\end{aligned}$$

- If  $G_m$  is a strong classifier overall, then its  $\alpha_m$  will be large; this means that if  $x_i$  is misclassified,  $w_i$  will increase to a greater extent

# AdaBoost: Algorithm

Given training set  $\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}$ .

- 1 Initialize observation weights  $w_i = 1, i = 1, 2, \dots, n$ .

# AdaBoost: Algorithm

Given training set  $\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}$ .

- ① Initialize observation weights  $w_i = 1, i = 1, 2, \dots, n$ .
- ② For  $m = 1$  to  $M$ :
  - ① Base learner fits weighted training data and returns  $G_m(x)$



# AdaBoost: Algorithm

Given training set  $\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}$ .

- ① Initialize observation weights  $w_i = 1, i = 1, 2, \dots, n$ .
- ② For  $m = 1$  to  $M$ :
  - ① Base learner fits weighted training data and returns  $G_m(x)$
  - ② Compute *weighted empirical 0-1 risk*:

$$\text{err}_m = \frac{1}{W} \sum_{i=1}^n w_i \mathbb{1}[y_i \neq G_m(x_i)] \quad \text{where } W = \sum_{i=1}^n w_i.$$

# AdaBoost: Algorithm

Given training set  $\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}$ .

- ① Initialize observation weights  $w_i = 1, i = 1, 2, \dots, n$ .
- ② For  $m = 1$  to  $M$ :
  - ① Base learner fits weighted training data and returns  $G_m(x)$
  - ② Compute *weighted empirical 0-1 risk*:

$$\text{err}_m = \frac{1}{W} \sum_{i=1}^n w_i \mathbb{1}[y_i \neq G_m(x_i)] \quad \text{where } W = \sum_{i=1}^n w_i.$$

- ③ Compute *classifier weight*:  $\alpha_m = \ln \left( \frac{1 - \text{err}_m}{\text{err}_m} \right)$ .

# AdaBoost: Algorithm

Given training set  $\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}$ .

- ① Initialize observation weights  $w_i = 1, i = 1, 2, \dots, n$ .
- ② For  $m = 1$  to  $M$ :
  - ① Base learner fits weighted training data and returns  $G_m(x)$
  - ② Compute *weighted empirical 0-1 risk*:

$$\text{err}_m = \frac{1}{W} \sum_{i=1}^n w_i \mathbb{1}[y_i \neq G_m(x_i)] \quad \text{where } W = \sum_{i=1}^n w_i.$$

- ③ Compute *classifier weight*:  $\alpha_m = \ln \left( \frac{1 - \text{err}_m}{\text{err}_m} \right)$ .
  - ④ Update *example weight*:  $w_i \leftarrow w_i \cdot \exp[\alpha_m \mathbb{1}[y_i \neq G_m(x_i)]]$

# AdaBoost: Algorithm

Given training set  $\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}$ .

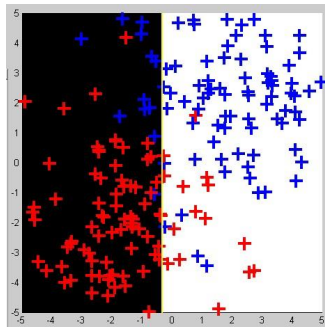
- ① Initialize observation weights  $w_i = 1, i = 1, 2, \dots, n$ .
- ② For  $m = 1$  to  $M$ :
  - ① Base learner fits weighted training data and returns  $G_m(x)$
  - ② Compute *weighted empirical 0-1 risk*:

$$\text{err}_m = \frac{1}{W} \sum_{i=1}^n w_i \mathbb{1}[y_i \neq G_m(x_i)] \quad \text{where } W = \sum_{i=1}^n w_i.$$

- ③ Compute *classifier weight*:  $\alpha_m = \ln \left( \frac{1 - \text{err}_m}{\text{err}_m} \right)$ .
  - ④ Update *example weight*:  $w_i \leftarrow w_i \cdot \exp[\alpha_m \mathbb{1}[y_i \neq G_m(x_i)]]$
- ③ Return *voted classifier*:  $G(x) = \text{sign} \left[ \sum_{m=1}^M \alpha_m G_m(x) \right]$ .

# AdaBoost with Decision Stumps


- After 1 round:



**Figure:** Size of plus sign represents weight of example. Blackness represents preference for red class; whiteness represents preference for blue class.

# AdaBoost with Decision Stumps


- After 3 rounds:



**Figure:** Size of plus sign represents weight of example. Blackness represents preference for red class; whiteness represents preference for blue class.

# AdaBoost with Decision Stumps

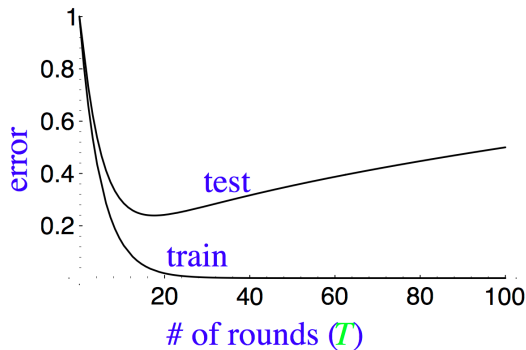
- After 120 rounds:



**Figure:** Size of plus sign represents weight of example. Blackness represents preference for red class; whiteness represents preference for blue class.

# Does AdaBoost overfit?

- Does a large number of rounds of boosting lead to overfitting?
- If we were overfitting, the learning curves would look like:




From Rob Schapire's NIPS 2007 Boosting tutorial.



# Learning Curves for AdaBoost

- AdaBoost is usually quite resistant to overfitting
- The test error continues to decrease even after the training error drops to zero!



From Rob Schapire's NIPS 2007 Boosting tutorial.

# AdaBoost for Face Detection


- Famous application of boosting: detecting faces in images (Viola & Jones, 2001)
- A few twists on standard algorithm

# AdaBoost for Face Detection

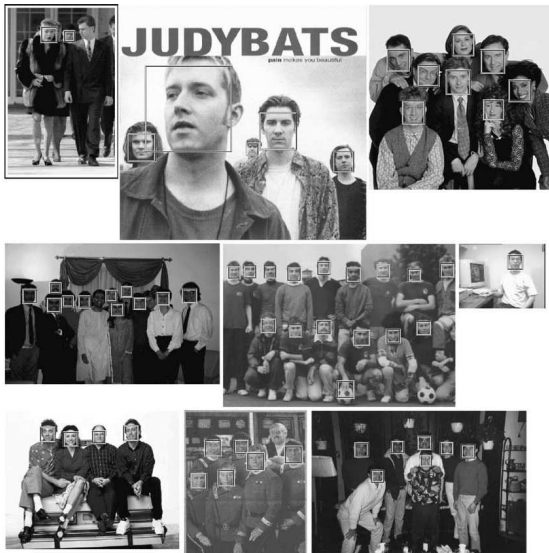
- Famous application of boosting: detecting faces in images (Viola & Jones, 2001)
- A few twists on standard algorithm
  - Pre-define weak classifiers, so optimization=selection

# AdaBoost for Face Detection

- Famous application of boosting: detecting faces in images (Viola & Jones, 2001)
- A few twists on standard algorithm
  - Pre-define weak classifiers, so optimization=selection
  - Smart way to do inference in real-time (in 2001 hardware)



# AdaBoost Face Detection Results



- Boosting is used to reduce bias from shallow decision trees

# Interim Summary

- Boosting is used to reduce bias from shallow decision trees
- Each classifier is trained to reduce errors of its previous ensemble.

## Interim Summary

- Boosting is used to reduce bias from shallow decision trees
- Each classifier is trained to reduce errors of its previous ensemble.
- AdaBoost is a very powerful off-the-shelf classifier.



# Interim Summary

- Boosting is used to reduce bias from shallow decision trees
- Each classifier is trained to reduce errors of its previous ensemble.
- AdaBoost is a very powerful off-the-shelf classifier.
- Next

# Interim Summary

- Boosting is used to reduce bias from shallow decision trees
- Each classifier is trained to reduce errors of its previous ensemble.
- AdaBoost is a very powerful off-the-shelf classifier.
- Next
  - What is the objective function of AdaBoost?

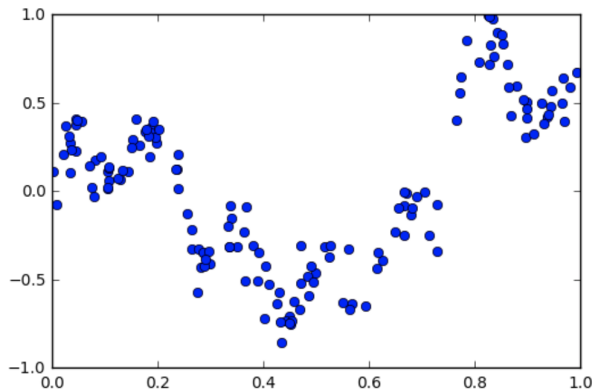
# Interim Summary

- Boosting is used to reduce bias from shallow decision trees
- Each classifier is trained to reduce errors of its previous ensemble.
- AdaBoost is a very powerful off-the-shelf classifier.
- Next
  - What is the objective function of AdaBoost?
  - Generalizations to other loss functions

- Boosting is used to reduce bias from shallow decision trees
- Each classifier is trained to reduce errors of its previous ensemble.
- AdaBoost is a very powerful off-the-shelf classifier.
- Next
  - What is the objective function of AdaBoost?
  - Generalizations to other loss functions
  - Gradient Boosting

# Nonlinear Regression

- How do we fit the following data?
- Another way to get non-linear models in a linear form—adaptive basis function models.



# Linear Model with Basis Functions

- Fit a linear combination of transformations of the input:

$$f(x) = \sum_{m=1}^M v_m h_m(x),$$

where  $h_m$ 's are called **basis functions** (or feature functions in ML):

$$h_1, \dots, h_M : \mathcal{X} \rightarrow \mathbb{R}$$

# Linear Model with Basis Functions

- Fit a linear combination of transformations of the input:

$$f(x) = \sum_{m=1}^M v_m h_m(x),$$

where  $h_m$ 's are called **basis functions** (or feature functions in ML):

$$h_1, \dots, h_M : \mathcal{X} \rightarrow \mathbb{R}$$

- Example: polynomial regression where  $h_m(x) = x^m$ .

# Linear Model with Basis Functions

- Fit a linear combination of transformations of the input:

$$f(x) = \sum_{m=1}^M v_m h_m(x),$$

where  $h_m$ 's are called **basis functions** (or feature functions in ML):

$$h_1, \dots, h_M : \mathcal{X} \rightarrow \mathbb{R}$$

- Example: polynomial regression where  $h_m(x) = x^m$ .
- Can we use this model for classification?



# Linear Model with Basis Functions

- Fit a linear combination of transformations of the input:

$$f(x) = \sum_{m=1}^M v_m h_m(x),$$

where  $h_m$ 's are called **basis functions** (or feature functions in ML):

$$h_1, \dots, h_M : \mathcal{X} \rightarrow \mathbb{R}$$

- Example: polynomial regression where  $h_m(x) = x^m$ .
- Can we use this model for classification?
- Can fit this using standard methods for linear models (e.g. least squares, lasso, ridge, etc.)

# Linear Model with Basis Functions

- Fit a linear combination of transformations of the input:

$$f(x) = \sum_{m=1}^M v_m h_m(x),$$

where  $h_m$ 's are called **basis functions** (or feature functions in ML):

$$h_1, \dots, h_M : \mathcal{X} \rightarrow \mathbb{R}$$

- Example: polynomial regression where  $h_m(x) = x^m$ .
- Can we use this model for classification?
- Can fit this using standard methods for linear models (e.g. least squares, lasso, ridge, etc.)
  - *Note that  $h_m$ 's are fixed and known, i.e. chosen ahead of time.*

# Adaptive Basis Function Model

- What if we want to learn the basis functions? (hence *adaptive*)

# Adaptive Basis Function Model

- What if we want to learn the basis functions? (hence *adaptive*)
- Base hypothesis space  $\mathcal{H}$  consisting of functions  $h : \mathcal{X} \rightarrow \mathbb{R}$ .
- An **adaptive basis function expansion** over  $\mathcal{H}$  is an ensemble model:

$$f(x) = \sum_{m=1}^M v_m h_m(x), \quad (4)$$

where  $v_m \in \mathbb{R}$  and  $h_m \in \mathcal{H}$ .

# Adaptive Basis Function Model

- What if we want to learn the basis functions? (hence *adaptive*)
- Base hypothesis space  $\mathcal{H}$  consisting of functions  $h : \mathcal{X} \rightarrow \mathbb{R}$ .
- An **adaptive basis function expansion** over  $\mathcal{H}$  is an ensemble model:

$$f(x) = \sum_{m=1}^M v_m h_m(x), \quad (4)$$

where  $v_m \in \mathbb{R}$  and  $h_m \in \mathcal{H}$ .

- Combined hypothesis space:

$$\mathcal{F}_M = \left\{ \sum_{m=1}^M v_m h_m(x) \mid v_m \in \mathbb{R}, h_m \in \mathcal{H}, m = 1, \dots, M \right\}$$

- What are the learnable?

# Empirical Risk Minimization

- What's our learning objective?

$$\hat{f} = \arg \min_{f \in \mathcal{F}_M} \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i)),$$

for some loss function  $\ell$ .

# Empirical Risk Minimization

- What's our learning objective?

$$\hat{f} = \arg \min_{f \in \mathcal{F}_M} \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i)),$$

for some loss function  $\ell$ .

- Write ERM objective function as

$$J(v_1, \dots, v_M, h_1, \dots, h_M) = \frac{1}{n} \sum_{i=1}^n \ell \left( y_i, \sum_{m=1}^M v_m h_m(x) \right).$$

# Empirical Risk Minimization

- What's our learning objective?

$$\hat{f} = \arg \min_{f \in \mathcal{F}_M} \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i)),$$

for some loss function  $\ell$ .

- Write ERM objective function as

$$J(v_1, \dots, v_M, h_1, \dots, h_M) = \frac{1}{n} \sum_{i=1}^n \ell \left( y_i, \sum_{m=1}^M v_m h_m(x) \right).$$

- How to optimize  $J$ ? i.e. how to learn?



- Suppose our base hypothesis space is parameterized by  $\Theta = \mathbb{R}^b$ :

$$J(v_1, \dots, v_M, \theta_1, \dots, \theta_M) = \frac{1}{n} \sum_{i=1}^n \ell \left( y_i, \sum_{m=1}^M v_m h(x; \theta_m) \right).$$

- Suppose our base hypothesis space is parameterized by  $\Theta = \mathbb{R}^b$ :

$$J(v_1, \dots, v_M, \theta_1, \dots, \theta_M) = \frac{1}{n} \sum_{i=1}^n \ell \left( y_i, \sum_{m=1}^M v_m h(x; \theta_m) \right).$$

- Can we optimize it with SGD?

- Suppose our base hypothesis space is parameterized by  $\Theta = \mathbb{R}^b$ :

$$J(v_1, \dots, v_M, \theta_1, \dots, \theta_M) = \frac{1}{n} \sum_{i=1}^n \ell \left( y_i, \sum_{m=1}^M v_m h(x; \theta_m) \right).$$

- Can we optimize it with SGD?
  - Can we differentiate  $J$  w.r.t.  $v_m$ 's and  $\theta_m$ 's?

# Gradient-Based Methods

- Suppose our base hypothesis space is parameterized by  $\Theta = \mathbb{R}^b$ :

$$J(v_1, \dots, v_M, \theta_1, \dots, \theta_M) = \frac{1}{n} \sum_{i=1}^n \ell \left( y_i, \sum_{m=1}^M v_m h(x; \theta_m) \right).$$

- Can we optimize it with SGD?
  - Can we differentiate  $J$  w.r.t.  $v_m$ 's and  $\theta_m$ 's?
- For some hypothesis spaces and typical loss functions, yes!
  - Neural networks fall into this category! ( $h_1, \dots, h_M$  are neurons of last hidden layer.)

# What if Gradient Based Methods Don't Apply?

What if base hypothesis space  $\mathcal{H}$  consists of decision trees?

# What if Gradient Based Methods Don't Apply?

What if base hypothesis space  $\mathcal{H}$  consists of decision trees?

- Can we even parameterize trees with  $\Theta = \mathbb{R}^b$ ?
- Even if we could, predictions would not change continuously w.r.t.  $\theta \in \Theta$ , so certainly not differentiable.

# What if Gradient Based Methods Don't Apply?

What if base hypothesis space  $\mathcal{H}$  consists of decision trees?

- Can we even parameterize trees with  $\Theta = \mathbb{R}^b$ ?
- Even if we could, predictions would not change continuously w.r.t.  $\theta \in \Theta$ , so certainly not differentiable.

What about a greedy algorithm similar to Adaboost?

# What if Gradient Based Methods Don't Apply?

What if base hypothesis space  $\mathcal{H}$  consists of decision trees?

- Can we even parameterize trees with  $\Theta = \mathbb{R}^b$ ?
- Even if we could, predictions would not change continuously w.r.t.  $\theta \in \Theta$ , so certainly not differentiable.

What about a greedy algorithm similar to Adaboost?

- Applies to non-parametric or non-differentiable basis functions.
- But is it optimizing our objective using some loss function?



Today we'll discuss **gradient boosting**.

- Gradient descent in the *function space*.
- It applies whenever
  - our loss function is [sub]differentiable w.r.t. training predictions  $f(x_i)$ , and
  - we can do regression with the base hypothesis space  $\mathcal{H}$ .

## Forward Stagewise Additive Modeling

# Forward Stagewise Additive Modeling (FSAM)

**Goal** fit model  $f(x) = \sum_{m=1}^M v_m h_m(x)$  given some loss function.

**Approach** Greedily fit one function at a time without adjusting previous functions, hence “forward stagewise”.

- After  $m-1$  stages, we have

$$f_{m-1} = \sum_{i=1}^{m-1} v_i h_i.$$

# Forward Stagewise Additive Modeling (FSAM)

**Goal** fit model  $f(x) = \sum_{m=1}^M v_m h_m(x)$  given some loss function.

**Approach** Greedily fit one function at a time without adjusting previous functions, hence “forward stagewise”.

- After  $m-1$  stages, we have

$$f_{m-1} = \sum_{i=1}^{m-1} v_i h_i.$$

- In  $m$ 'th round, we want to find  $h_m \in \mathcal{H}$  (i.e. a basis function) and  $v_m > 0$  such that

$$f_m = \underbrace{f_{m-1}}_{\text{fixed}} + v_m h_m$$

improves objective function value by as much as possible.

# Forward Stagewise Additive Modeling for ERM

Let's plug in our objective function.

- 1 Initialize  $f_0(x) = 0$ .
- 2 For  $m = 1$  to  $M$ :

# Forward Stagewise Additive Modeling for ERM

Let's plug in our objective function.

- 1 Initialize  $f_0(x) = 0$ .
- 2 For  $m = 1$  to  $M$ :
  - 1 Compute:

$$(v_m, h_m) = \arg \min_{v \in \mathbb{R}, h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \ell \left( y_i, f_{m-1}(x_i) + \underbrace{vh(x_i)}_{\text{new piece}} \right).$$

# Forward Stagewise Additive Modeling for ERM

Let's plug in our objective function.

❶ Initialize  $f_0(x) = 0$ .

❷ For  $m = 1$  to  $M$ :

❶ Compute:

$$(v_m, h_m) = \arg \min_{v \in \mathbb{R}, h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \ell \left( y_i, f_{m-1}(x_i) + \underbrace{v h(x_i)}_{\text{new piece}} \right).$$

❷ Set  $f_m = f_{m-1} + v_m h_m$ .

# Forward Stagewise Additive Modeling for ERM

Let's plug in our objective function.

❶ Initialize  $f_0(x) = 0$ .

❷ For  $m = 1$  to  $M$ :

❶ Compute:

$$(v_m, h_m) = \arg \min_{v \in \mathbb{R}, h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \ell \left( y_i, f_{m-1}(x_i) + \underbrace{vh(x_i)}_{\text{new piece}} \right).$$

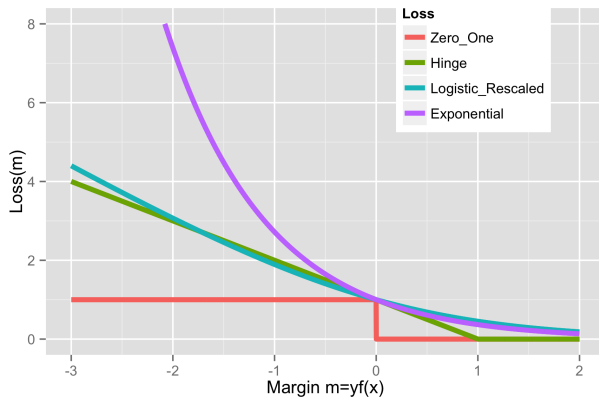
❷ Set  $f_m = f_{m-1} + v_m h_m$ .

❸ Return:  $f_M$ .



# Exponential Loss

- Introduce the **exponential loss**:  $\ell(y, f(x)) = \exp\left(\underbrace{-yf(x)}_{\text{margin}}\right)$ .



# Forward Stagewise Additive Modeling with exponential loss

Recall that we want to do FSAM with exponential loss.

❶ Initialize  $f_0(x) = 0$ .

❷ For  $m = 1$  to  $M$ :

❶ Compute:

$$(v_m, h_m) = \arg \min_{v \in \mathbb{R}, h \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n \ell_{\text{exp}} \left( y_i, f_{m-1}(x_i) + \underbrace{vh(x_i)}_{\text{new piece}} \right).$$

❷ Set  $f_m = f_{m-1} + v_m h_m$ .

❸ Return:  $f_M$ .

## FSAM with Exponential Loss: objective function

- Base hypothesis:  $\mathcal{H} = \{h: \mathcal{X} \rightarrow \{-1, 1\}\}$ .
- Objective function in the  $m$ 'th round:

## FSAM with Exponential Loss: objective function

- Base hypothesis:  $\mathcal{H} = \{h: \mathcal{X} \rightarrow \{-1, 1\}\}$ .
- Objective function in the  $m$ 'th round:

$$J(v, h) = \sum_{i=1}^n \exp[-y_i (f_{m-1}(x_i) + v h(x_i))] \quad (5)$$

$$= \sum_{i=1}^n w_i^m \exp[-y_i v h(x_i)] \quad w_i^m \stackrel{\text{def}}{=} \exp[-y_i f_{m-1}(x_i)] \quad (6)$$

$$= \sum_{i=1}^n w_i^m [\mathbb{I}(y_i = h(x_i)) e^{-v} + \mathbb{I}(y_i \neq h(x_i)) e^v] \quad h(x_i) \in \{1, -1\} \quad (7)$$

$$= \sum_{i=1}^n w_i^m [(e^v - e^{-v}) \mathbb{I}(y_i \neq h(x_i)) + e^{-v}] \quad \mathbb{I}(y_i = h(x_i)) = 1 - \mathbb{I}(y_i \neq h(x_i)) \quad (8)$$

## FSAM with Exponential Loss: basis function

- Objective function in the  $m$ 'th round:

$$J(v, h) = \sum_{i=1}^n w_i^m \left[ (e^v - e^{-v}) \mathbb{I}(y_i \neq h(x_i)) + e^{-v} \right]. \quad (9)$$

## FSAM with Exponential Loss: basis function

- Objective function in the  $m$ 'th round:

$$J(v, h) = \sum_{i=1}^n w_i^m [(e^v - e^{-v})\mathbb{I}(y_i \neq h(x_i)) + e^{-v}]. \quad (9)$$

- If  $v > 0$ , then

$$\arg \min_{h \in \mathcal{H}} J(v, h) = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n w_i^m \mathbb{I}(y_i \neq h(x_i)) \quad (10)$$

(12)

## FSAM with Exponential Loss: basis function

- Objective function in the  $m$ 'th round:

$$J(v, h) = \sum_{i=1}^n w_i^m [(e^v - e^{-v})\mathbb{I}(y_i \neq h(x_i)) + e^{-v}]. \quad (9)$$

- If  $v > 0$ , then

$$\arg \min_{h \in \mathcal{H}} J(v, h) = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n w_i^m \mathbb{I}(y_i \neq h(x_i)) \quad (10)$$

$$h_m = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n w_i^m \mathbb{I}(y_i \neq h(x_i)) \quad (11)$$

(12)

## FSAM with Exponential Loss: basis function

- Objective function in the  $m$ 'th round:

$$J(v, h) = \sum_{i=1}^n w_i^m [(e^v - e^{-v})\mathbb{I}(y_i \neq h(x_i)) + e^{-v}]. \quad (9)$$

- If  $v > 0$ , then

$$\arg \min_{h \in \mathcal{H}} J(v, h) = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n w_i^m \mathbb{I}(y_i \neq h(x_i)) \quad (10)$$

$$h_m = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n w_i^m \mathbb{I}(y_i \neq h(x_i)) \quad (11)$$

$$= \arg \min_{h \in \mathcal{H}} \frac{1}{\sum_{i=1}^n w_i^m} \sum_{i=1}^n w_i^m \mathbb{I}(y_i \neq h(x_i)) \quad \text{multiply by a positive constant} \quad (12)$$



## FSAM with Exponential Loss: basis function

- Objective function in the  $m$ 'th round:

$$J(v, h) = \sum_{i=1}^n w_i^m [(e^v - e^{-v})\mathbb{I}(y_i \neq h(x_i)) + e^{-v}]. \quad (9)$$

- If  $v > 0$ , then

$$\arg \min_{h \in \mathcal{H}} J(v, h) = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n w_i^m \mathbb{I}(y_i \neq h(x_i)) \quad (10)$$

$$h_m = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n w_i^m \mathbb{I}(y_i \neq h(x_i)) \quad (11)$$

$$= \arg \min_{h \in \mathcal{H}} \frac{1}{\sum_{i=1}^n w_i^m} \sum_{i=1}^n w_i^m \mathbb{I}(y_i \neq h(x_i)) \quad \text{multiply by a positive constant} \quad (12)$$

i.e.  $h_m$  is the minimizer of the weighted zero-one loss.

## FSAM with Exponential Loss: classifier weights

- Define the weighted zero-one error:

$$\text{err}_m = \frac{\sum_{i=1}^n w_i^m \mathbb{I}(y_i \neq h(x_i))}{\sum_{i=1}^n w_i^m}. \quad (13)$$

## FSAM with Exponential Loss: classifier weights

- Define the weighted zero-one error:

$$\text{err}_m = \frac{\sum_{i=1}^n w_i^m \mathbb{I}(y_i \neq h(x_i))}{\sum_{i=1}^n w_i^m}. \quad (13)$$

- Exercise:** show that the optimal  $v$  is:

$$v_m = \frac{1}{2} \log \frac{1 - \text{err}_m}{\text{err}_m} \quad (14)$$

## FSAM with Exponential Loss: classifier weights

- Define the weighted zero-one error:

$$\text{err}_m = \frac{\sum_{i=1}^n w_i^m \mathbb{I}(y_i \neq h(x_i))}{\sum_{i=1}^n w_i^m}. \quad (13)$$

- Exercise:** show that the optimal  $v$  is:

$$v_m = \frac{1}{2} \log \frac{1 - \text{err}_m}{\text{err}_m} \quad (14)$$

- Same as the classifier weights in Adaboost (differ by a constant).

## FSAM with Exponential Loss: classifier weights

- Define the weighted zero-one error:

$$\text{err}_m = \frac{\sum_{i=1}^n w_i^m \mathbb{I}(y_i \neq h(x_i))}{\sum_{i=1}^n w_i^m}. \quad (13)$$

- Exercise:** show that the optimal  $v$  is:

$$v_m = \frac{1}{2} \log \frac{1 - \text{err}_m}{\text{err}_m} \quad (14)$$

- Same as the classifier weights in Adaboost (differ by a constant).
- If  $\text{err}_m < 0.5$  (better than chance), then  $v_m > 0$ .

## FSAM with Exponential Loss: example weights

- Weights in the next round:

$$w_i^{m+1} \stackrel{\text{def}}{=} \exp[-y_i f_m(x_i)] \quad (15)$$

(18)

# FSAM with Exponential Loss: example weights

- Weights in the next round:

$$w_i^{m+1} \stackrel{\text{def}}{=} \exp[-y_i f_m(x_i)] \quad (15)$$

$$= w_i^m \exp[-y_i v_m h_m(x_i)] \quad f_m(x_i) = f_{m-1}(x_i) + v_m h_m(x_i) \quad (16)$$

$$= w_i^m \exp[-v_m \mathbb{I}(y_i = h_m(x_i)) + v_m \mathbb{I}(y_i \neq h_m(x_i))] \quad (17)$$

$$= w_i^m \exp[2v_m \mathbb{I}(y_i \neq h_m(x_i))] \underbrace{\exp^{-v_m}}_{\text{scaler}} \quad (18)$$

## FSAM with Exponential Loss: example weights

- Weights in the next round:

$$w_i^{m+1} \stackrel{\text{def}}{=} \exp[-y_i f_m(x_i)] \quad (15)$$

$$= w_i^m \exp[-y_i v_m h_m(x_i)] \quad f_m(x_i) = f_{m-1}(x_i) + v_m h_m(x_i) \quad (16)$$

$$= w_i^m \exp[-v_m \mathbb{I}(y_i = h_m(x_i)) + v_m \mathbb{I}(y_i \neq h_m(x_i))] \quad (17)$$

$$= w_i^m \exp[2v_m \mathbb{I}(y_i \neq h_m(x_i))] \underbrace{\exp^{-v_m}}_{\text{scaler}} \quad (18)$$

- The constant scaler will cancel out during normalization.



## FSAM with Exponential Loss: example weights

- Weights in the next round:

$$w_i^{m+1} \stackrel{\text{def}}{=} \exp[-y_i f_m(x_i)] \quad (15)$$

$$= w_i^m \exp[-y_i v_m h_m(x_i)] \quad f_m(x_i) = f_{m-1}(x_i) + v_m h_m(x_i) \quad (16)$$

$$= w_i^m \exp[-v_m \mathbb{I}(y_i = h_m(x_i)) + v_m \mathbb{I}(y_i \neq h_m(x_i))] \quad (17)$$

$$= w_i^m \exp[2v_m \mathbb{I}(y_i \neq h_m(x_i))] \underbrace{\exp^{-v_m}}_{\text{scaler}} \quad (18)$$

- The constant scaler will cancel out during normalization.
- $2v_m = \alpha_m$  in Adaboost.

# Why Exponential Loss

- $\ell_{\text{exp}}(y, f(x)) = \exp(-yf(x))$ .

# Why Exponential Loss

- $\ell_{\text{exp}}(y, f(x)) = \exp(-yf(x))$ .
- **Exercise:** show that the optimal estimate is

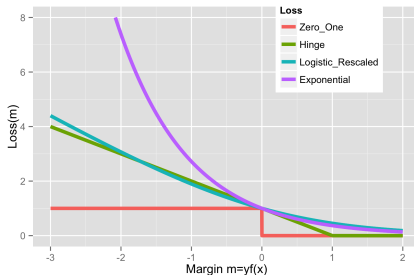
$$f^*(x) = \frac{1}{2} \log \frac{p(y=1|x)}{p(y=0|x)}. \quad (19)$$

# Why Exponential Loss

- $\ell_{\text{exp}}(y, f(x)) = \exp(-yf(x))$ .
- **Exercise:** show that the optimal estimate is

$$f^*(x) = \frac{1}{2} \log \frac{p(y = 1 | x)}{p(y = 0 | x)}. \quad (19)$$

- How is it different from other losses?



- Exponential loss puts a high penalty on misclassified examples.

## AdaBoost / Exponential Loss: Robustness Issues

- Exponential loss puts a high penalty on misclassified examples.
  - $\implies$  not robust to outliers / noise.
- Empirically, AdaBoost has degraded performance in situations with
  - high Bayes error rate (intrinsic randomness in the label)

## AdaBoost / Exponential Loss: Robustness Issues

- Exponential loss puts a high penalty on misclassified examples.
  - $\implies$  not robust to outliers / noise.
- Empirically, AdaBoost has degraded performance in situations with
  - high Bayes error rate (intrinsic randomness in the label)
- Logistic/Log loss performs better in settings with high Bayes error.
- Exponential loss has some computational advantages over log loss though.

We've seen

- Use basis function to obtain *nonlinear* models:  $f(x) = \sum_{i=1}^M v_m h_m(x)$  with known  $h_m$ 's.
- *Adaptive* basis function models:  $f(x) = \sum_{i=1}^M v_m h_m(x)$  with unknown  $h_m$ 's.
- Forward stagewise additive modeling: greedily fit  $h_m$ 's to minimize the average loss.



We've seen

- Use basis function to obtain *nonlinear* models:  $f(x) = \sum_{i=1}^M v_m h_m(x)$  with known  $h_m$ 's.
- *Adaptive* basis function models:  $f(x) = \sum_{i=1}^M v_m h_m(x)$  with unknown  $h_m$ 's.
- Forward stagewise additive modeling: greedily fit  $h_m$ 's to minimize the average loss.

But,

- We only know how to do FSAM for certain loss functions.
- Need to derive new algorithms for different loss functions.

Next, how to do FSAM in general.

## Gradient Boosting / “Anyboost”

## FSAM with squared loss

- Objective function at  $m$ 'th round:

$$J(v, h) = \frac{1}{n} \sum_{i=1}^n \left( y_i - \left[ f_{m-1}(x_i) + \underbrace{vh(x_i)}_{\text{new piece}} \right] \right)^2$$

## FSAM with squared loss

- Objective function at  $m$ 'th round:

$$J(v, h) = \frac{1}{n} \sum_{i=1}^n \left( y_i - \left[ f_{m-1}(x_i) + \underbrace{vh(x_i)}_{\text{new piece}} \right] \right)^2$$

- If  $\mathcal{H}$  is closed under rescaling (i.e. if  $h \in \mathcal{H}$ , then  $vh \in \mathcal{H}$  for all  $h \in \mathbb{R}$ ), then don't need  $v$ .

## FSAM with squared loss

- Objective function at  $m$ 'th round:

$$J(v, h) = \frac{1}{n} \sum_{i=1}^n \left( y_i - \left[ f_{m-1}(x_i) + \underbrace{vh(x_i)}_{\text{new piece}} \right] \right)^2$$

- If  $\mathcal{H}$  is closed under rescaling (i.e. if  $h \in \mathcal{H}$ , then  $vh \in \mathcal{H}$  for all  $h \in \mathcal{H}$ ), then don't need  $v$ .
- Take  $v = 1$  and minimize

$$J(h) = \frac{1}{n} \sum_{i=1}^n \left( \left[ \underbrace{y_i - f_{m-1}(x_i)}_{\text{residual}} \right] - h(x_i) \right)^2$$

## FSAM with squared loss

- Objective function at  $m$ 'th round:

$$J(v, h) = \frac{1}{n} \sum_{i=1}^n \left( y_i - \left[ f_{m-1}(x_i) + \underbrace{vh(x_i)}_{\text{new piece}} \right] \right)^2$$

- If  $\mathcal{H}$  is closed under rescaling (i.e. if  $h \in \mathcal{H}$ , then  $vh \in \mathcal{H}$  for all  $h \in \mathcal{H}$ ), then don't need  $v$ .
- Take  $v = 1$  and minimize

$$J(h) = \frac{1}{n} \sum_{i=1}^n \left( \left[ \underbrace{y_i - f_{m-1}(x_i)}_{\text{residual}} \right] - h(x_i) \right)^2$$

## FSAM with squared loss

- Objective function at  $m$ 'th round:

$$J(v, h) = \frac{1}{n} \sum_{i=1}^n \left( y_i - \left[ f_{m-1}(x_i) + \underbrace{vh(x_i)}_{\text{new piece}} \right] \right)^2$$

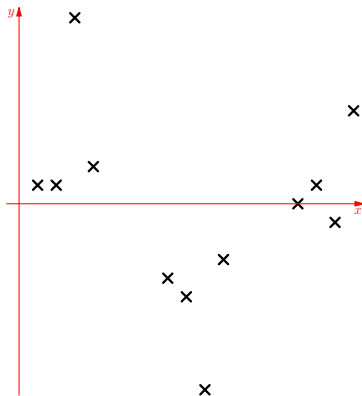
- If  $\mathcal{H}$  is closed under rescaling (i.e. if  $h \in \mathcal{H}$ , then  $vh \in \mathcal{H}$  for all  $v \in \mathbb{R}$ ), then don't need  $v$ .
- Take  $v = 1$  and minimize

$$J(h) = \frac{1}{n} \sum_{i=1}^n \left( \left[ \underbrace{y_i - f_{m-1}(x_i)}_{\text{residual}} \right] - h(x_i) \right)^2$$

- This is just fitting the residuals with least-squares regression!
- Example base hypothesis space: regression stumps.

# $L^2$ Boosting with Decision Stumps: Demo

- Consider FSAM with  $L^2$  loss (i.e.  $L^2$  Boosting)
- For base hypothesis space of **regression stumps**

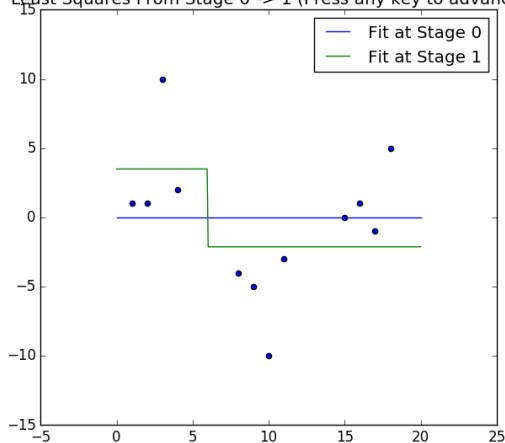


Plot courtesy of Brett Bernstein.

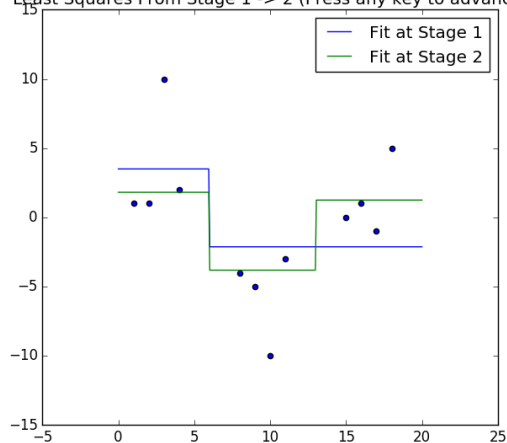


# $L^2$ Boosting with Decision Stumps: Results

Least Squares From Stage 0 -> 1 (Press any key to advance)

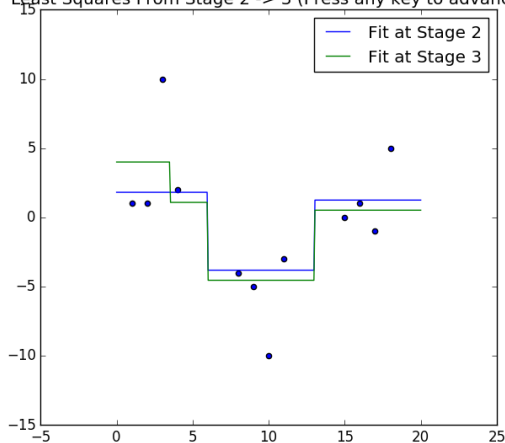


Least Squares From Stage 1 -> 2 (Press any key to advance)

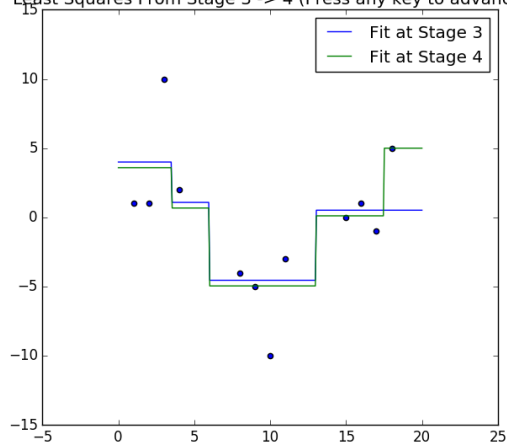


# $L^2$ Boosting with Decision Stumps: Results

Least Squares From Stage 2 -> 3 (Press any key to advance)

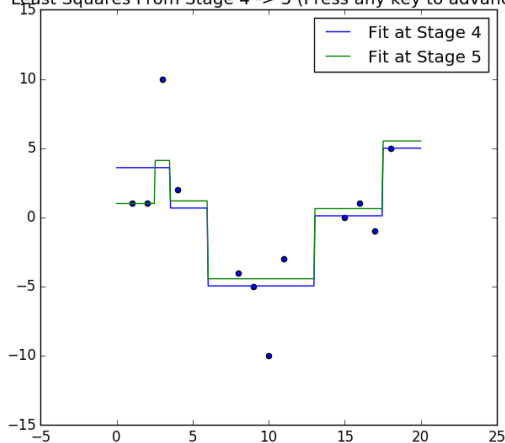


Least Squares From Stage 3 -> 4 (Press any key to advance)

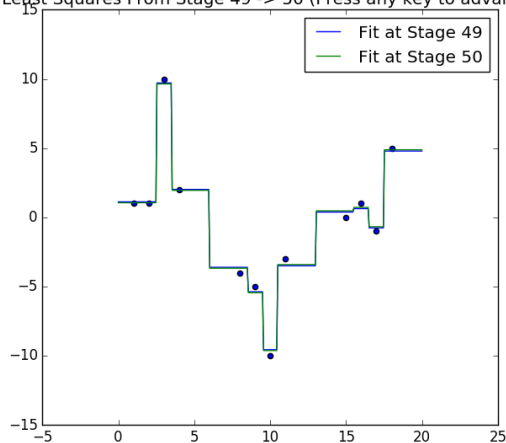


# $L^2$ Boosting with Decision Stumps: Results

Least Squares From Stage 4 -> 5 (Press any key to advance)



Least Squares From Stage 49 -> 50 (Press any key to advance)



## Interpret the residual

- Objective:  $J(f) = \frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2$ .

## Interpret the residual

- Objective:  $J(f) = \frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2$ .
- What is the residual at  $x = x_i$ ?

$$\frac{\partial}{\partial f(x_i)} J(f) = \quad (20)$$

## Interpret the residual

- Objective:  $J(f) = \frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2$ .
- What is the residual at  $x = x_i$ ?

$$\frac{\partial}{\partial f(x_i)} J(f) = -2(y_i - f(x_i)) \quad (20)$$

- Gradient w.r.t.  $f$ : how should the output of  $f$  change to minimize the squared loss.
- *Residual is the negative gradient* (differ by some constant).

## Interpret the residual

- Objective:  $J(f) = \frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2$ .
- What is the residual at  $x = x_i$ ?

$$\frac{\partial}{\partial f(x_i)} J(f) = -2(y_i - f(x_i)) \quad (20)$$

- Gradient w.r.t.  $f$ : how should the output of  $f$  change to minimize the squared loss.
- *Residual is the negative gradient* (differ by some constant).
- At each boosting round, we learn a function  $h \in \mathcal{H}$  to fit the residual.

$$f \leftarrow f + \textcolor{red}{v} \textcolor{green}{h} \quad \text{FSAM / boosting} \quad (21)$$

(22)

## Interpret the residual

- Objective:  $J(f) = \frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2$ .
- What is the residual at  $x = x_i$ ?

$$\frac{\partial}{\partial f(x_i)} J(f) = -2(y_i - f(x_i)) \quad (20)$$

- Gradient w.r.t.  $f$ : how should the output of  $f$  change to minimize the squared loss.
- *Residual is the negative gradient* (differ by some constant).
- At each boosting round, we learn a function  $h \in \mathcal{H}$  to fit the residual.

$$f \leftarrow f + \nu h \quad \text{FSAM / boosting} \quad (21)$$

$$f \leftarrow f - \alpha \nabla_f J(f) \quad \text{gradient descent} \quad (22)$$

- $h$  approximates the gradient (step direction),  $\nu$  is the step size.



# “Functional” Gradient Descent

- We want to minimize

$$J(f) = \sum_{i=1}^n \ell(y_i, f(x_i)).$$

- In some sense, we want to take the gradient w.r.t.  $f$ .

# “Functional” Gradient Descent

- We want to minimize

$$J(f) = \sum_{i=1}^n \ell(y_i, f(x_i)).$$

- In some sense, we want to take the gradient w.r.t.  $f$ .
- $J(f)$  only depends on  $f$  at the  $n$  training points.

# “Functional” Gradient Descent

- We want to minimize

$$J(f) = \sum_{i=1}^n \ell(y_i, f(x_i)).$$

- In some sense, we want to take the gradient w.r.t.  $f$ .
- $J(f)$  only depends on  $f$  at the  $n$  training points.
- Define “parameters”

$$\mathbf{f} = (f(x_1), \dots, f(x_n))^T$$

and write the objective function as

$$J(\mathbf{f}) = \sum_{i=1}^n \ell(y_i, \mathbf{f}_i).$$

# Functional Gradient Descent: Unconstrained Step Direction

- Consider gradient descent on

$$J(f) = \sum_{i=1}^n \ell(y_i, f_i).$$

# Functional Gradient Descent: Unconstrained Step Direction

- Consider gradient descent on

$$J(f) = \sum_{i=1}^n \ell(y_i, f_i).$$

- The negative gradient step direction at  $f$  is

$$\begin{aligned} -g &= -\nabla_f J(f) \\ &= -(\partial_{f_1} \ell(y_1, f_1), \dots, \partial_{f_n} \ell(y_n, f_n)) \end{aligned}$$

which we can easily calculate.

- $-g \in \mathbb{R}^n$  is the direction we want to change each of our  $n$  predictions on training data.

# Functional Gradient Descent: Unconstrained Step Direction

- Consider gradient descent on

$$J(f) = \sum_{i=1}^n \ell(y_i, f_i).$$

- The negative gradient step direction at  $f$  is

$$\begin{aligned} -g &= -\nabla_f J(f) \\ &= -(\partial_{f_1} \ell(y_1, f_1), \dots, \partial_{f_n} \ell(y_n, f_n)) \end{aligned}$$

which we can easily calculate.

- $-g \in \mathbb{R}^n$  is the direction we want to change each of our  $n$  predictions on training data.
- With gradient descent, our final predictor will be an additive model:  $f_0 + \sum_{m=1}^M v_t(-g_t)$ .

# Functional Gradient Descent: Projection Step

- Unconstrained step direction is

$$-g = -\nabla_{\mathbf{f}} J(\mathbf{f}) = -(\partial_{f_1} \ell(y_1, f_1), \dots, \partial_{f_n} \ell(y_n, f_n)).$$

- Also called the “**pseudo-residuals**”. (For squared loss, they’re exactly the residuals.)

# Functional Gradient Descent: Projection Step

- Unconstrained step direction is

$$-g = -\nabla_{\mathbf{f}} J(\mathbf{f}) = -(\partial_{\mathbf{f}_1} \ell(y_1, \mathbf{f}_1), \dots, \partial_{\mathbf{f}_n} \ell(y_n, \mathbf{f}_n)).$$

- Also called the “**pseudo-residuals**”. (For squared loss, they’re exactly the residuals.)
- **Problem**: only know how to update at  $n$  points. How do we take a gradient step in  $\mathcal{H}$ ?



# Functional Gradient Descent: Projection Step

- Unconstrained step direction is

$$-g = -\nabla_{\mathbf{f}} J(\mathbf{f}) = -(\partial_{f_1} \ell(y_1, f_1), \dots, \partial_{f_n} \ell(y_n, f_n)).$$

- Also called the “**pseudo-residuals**”. (For squared loss, they’re exactly the residuals.)
- **Problem**: only know how to update at  $n$  points. How do we take a gradient step in  $\mathcal{H}$ ?
- **Solution**: approximate by the closest base hypothesis  $h \in \mathcal{H}$  (in the  $\ell^2$  sense):

$$\min_{h \in \mathcal{H}} \sum_{i=1}^n (-g_i - h(x_i))^2. \quad \text{least square regression} \quad (23)$$

- Take the  $h \in \mathcal{H}$  that best approximates  $-g$  as our step direction.

## Recap

- Objective function:

$$J(f) = \sum_{i=1}^n \ell(y_i, f(x_i)). \quad (24)$$

## Recap

- Objective function:

$$J(\mathbf{f}) = \sum_{i=1}^n \ell(y_i, f(x_i)). \quad (24)$$

- Unconstrained gradient  $\mathbf{g} \in \mathbb{R}^n$  w.r.t.  $\mathbf{f} = (f(x_1), \dots, f(x_n))^T$ :

$$\mathbf{g} = \nabla_{\mathbf{f}} J(\mathbf{f}) = (\partial_{f_1} \ell(y_1, f_1), \dots, \partial_{f_n} \ell(y_n, f_n)). \quad (25)$$

# Recap

- Objective function:

$$J(\mathbf{f}) = \sum_{i=1}^n \ell(y_i, f(x_i)). \quad (24)$$

- Unconstrained gradient  $\mathbf{g} \in \mathbb{R}^n$  w.r.t.  $\mathbf{f} = (f(x_1), \dots, f(x_n))^T$ :

$$\mathbf{g} = \nabla_{\mathbf{f}} J(\mathbf{f}) = (\partial_{f_1} \ell(y_1, f_1), \dots, \partial_{f_n} \ell(y_n, f_n)). \quad (25)$$

- Projected negative gradient  $h \in \mathcal{H}$ :

$$h = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n (-g_i - h(x_i))^2. \quad (26)$$

## Recap

- Objective function:

$$J(f) = \sum_{i=1}^n \ell(y_i, f(x_i)). \quad (24)$$

- Unconstrained gradient  $\mathbf{g} \in \mathbb{R}^n$  w.r.t.  $\mathbf{f} = (f(x_1), \dots, f(x_n))^T$ :

$$\mathbf{g} = \nabla_{\mathbf{f}} J(\mathbf{f}) = (\partial_{f_1} \ell(y_1, f_1), \dots, \partial_{f_n} \ell(y_n, f_n)). \quad (25)$$

- Projected negative gradient  $h \in \mathcal{H}$ :

$$h = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n (-g_i - h(x_i))^2. \quad (26)$$

- Gradient descent:

$$f \leftarrow f + \textcolor{red}{\alpha} h \quad (27)$$

# Functional Gradient Descent: hyperparameters

- Choose a step size by **line search**.

$$v_m = \arg \min_v \sum_{i=1}^n \ell\{y_i, f_{m-1}(x_i) + v h_m(x_i)\}.$$

- Not necessary. Can also choose a fixed hyperparameter  $v$ .
- Regularization through **shrinkage**:

$$f_m \leftarrow f_{m-1} + \lambda v_m h_m \quad \text{where } \lambda \in [0, 1]. \quad (28)$$

- Typically choose  $\lambda = 0.1$ .
- Choose  $M$ , i.e. when to stop.
  - Tune on validation set.

# Gradient boosting algorithm

- 1 Initialize  $f$  to a constant:  $f_0(x) = \arg \min_{\gamma} \sum_{i=1}^n \ell(y_i, \gamma)$ .
- 2 For  $m$  from 1 to  $M$ :
  - 1 Compute the pseudo-residuals (negative gradient):

$$r_{im} = - \left[ \frac{\partial}{\partial f(x_i)} \ell(y_i, f(x_i)) \right]_{f(x_i) = f_{m-1}(x_i)} \quad (29)$$

- 2 Fit a base learner  $h_m$  with squared loss using the dataset  $\{(x_i, r_{im})\}_{i=1}^n$ .
  - 3 [Optional] Find the best step size  $v_m = \arg \min_v \sum_{i=1}^n \ell(y_i, f_{m-1}(x_i) + v h_m(x_i))$ .
  - 4 Update  $f_m = f_{m-1} + \lambda v_m h_m$
- 3 Return  $f_M(x)$ .

# The Gradient Boosting Machine Ingredients (Recap)

- Take any loss function [sub]differentiable w.r.t. the prediction  $f(x_i)$
- Choose a base hypothesis space for regression.
- Choose number of steps (or a stopping criterion).
- Choose step size methodology.
- Then you're good to go!



# BinomialBoost: Gradient Boosting with Logistic Loss

- Recall the logistic loss for classification, with  $\mathcal{Y} = \{-1, 1\}$ :

$$\ell(y, f(x)) = \log \left( 1 + e^{-yf(x)} \right)$$

# BinomialBoost: Gradient Boosting with Logistic Loss

- Recall the logistic loss for classification, with  $\mathcal{Y} = \{-1, 1\}$ :

$$\ell(y, f(x)) = \log(1 + e^{-yf(x)})$$

- Pseudoresidual for  $i$ 'th example is negative derivative of loss w.r.t. prediction:

$$r_i = -\frac{\partial}{\partial f(x_i)} \ell(y_i, f(x_i)) \quad (30)$$

$$= -\frac{\partial}{\partial f(x_i)} \left[ \log(1 + e^{-y_i f(x_i)}) \right] \quad (31)$$

$$= \frac{y_i e^{-y_i f(x_i)}}{1 + e^{-y_i f(x_i)}} \quad (32)$$

$$= \frac{y_i}{1 + e^{y_i f(x_i)}} \quad (33)$$

# BinomialBoost: Gradient Boosting with Logistic Loss

- Pseudoresidual for  $i$ th example:

$$r_i = -\frac{\partial}{\partial f(x_i)} \left[ \log \left( 1 + e^{-y_i f(x_i)} \right) \right] = \frac{y_i}{1 + e^{y_i f(x_i)}}$$

# BinomialBoost: Gradient Boosting with Logistic Loss

- Pseudoresidual for  $i$ th example:

$$r_i = -\frac{\partial}{\partial f(x_i)} \left[ \log \left( 1 + e^{-y_i f(x_i)} \right) \right] = \frac{y_i}{1 + e^{y_i f(x_i)}}$$

- So if  $f_{m-1}(x)$  is prediction after  $m-1$  rounds, step direction for  $m$ 'th round is

$$h_m = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n \left[ \left( \frac{y_i}{1 + e^{y_i f_{m-1}(x_i)}} \right) - h(x_i) \right]^2.$$

# BinomialBoost: Gradient Boosting with Logistic Loss

- Pseudoresidual for  $i$ th example:

$$r_i = -\frac{\partial}{\partial f(x_i)} \left[ \log \left( 1 + e^{-y_i f(x_i)} \right) \right] = \frac{y_i}{1 + e^{y_i f(x_i)}}$$

- So if  $f_{m-1}(x)$  is prediction after  $m-1$  rounds, step direction for  $m$ 'th round is

$$h_m = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n \left[ \left( \frac{y_i}{1 + e^{y_i f_{m-1}(x_i)}} \right) - h(x_i) \right]^2.$$

- And  $f_m(x) = f_{m-1}(x) + \eta h_m(x)$ .

# Gradient Tree Boosting


- One common form of gradient boosting machine takes

$$\mathcal{H} = \{\text{regression trees of size } S\},$$

where  $S$  is the number of terminal nodes.


- $S = 2$  gives decision stumps
- Common choice:  $4 \leq S \leq 8$
- Software packages:
  - Gradient tree boosting is implemented by the `gbm` package for R
  - as `GradientBoostingClassifier` and `GradientBoostingRegressor` in `sklearn`
  - `xgboost` and `lightGBM` are state of the art for speed and performance

# Sinc Function: Our Dataset



From Natekin and Knoll's "Gradient boosting machines, a tutorial"

# Minimizing Square Loss with Ensemble of Decision Stumps



Decision stumps with 1, 10, 50, and 100 steps, shrinkage  $\lambda = 1$ .

Figure 3 from Natekin and Knoll's "Gradient boosting machines, a tutorial"




# Gradient Boosting in Practice

# Prevent overfitting

- Boosting is resistant to overfitting. Some explanations:
  - Implicit feature selection: greedily selects the best feature (weak learner)
  - As training goes on, impact of change is localized.
- But it can of course overfit. Common regularization methods:
  - Shrinkage (small learning rate)
  - Stochastic gradient boosting (row subsampling)
  - Feature subsampling (column subsampling)

# Step Size as Regularization



- (continued) sinc function regression
- Performance vs rounds of boosting and shrinkage. (Left is training set, right is validation set)

Figure 5 from Natekin and Knoll's "Gradient boosting machines, a tutorial"

# Stochastic Gradient Boosting

- For each stage,
  - choose random *subset of data* for computing projected gradient step.

# Stochastic Gradient Boosting

- For each stage,
  - choose random *subset of data* for computing projected gradient step.
- Why do this?
  - Introduce randomization thus may help overfitting.
  - Faster; often better than gradient descent given the same computation resource.

---

Introduced by Friedman (1999) in [Stochastic Gradient Boosting](#).

# Stochastic Gradient Boosting

- For each stage,
  - choose random *subset of data* for computing projected gradient step.
- Why do this?
  - Introduce randomization thus may help overfitting.
  - Faster; often better than gradient descent given the same computation resource.
- We can view this is a **minibatch method**.
  - Estimate the “true” step direction using a subset of data.

## Column / Feature Subsampling

- Similar to random forest, randomly choose *a subset of features* for each round.
- XGBoost paper says: “According to user feedback, using column sub-sampling prevents overfitting even more so than the traditional row sub-sampling.”
- Speeds up computation.

# Summary

- Motivating idea of boosting: combine weak learners to produce a strong learner.
- The statistical view: boosting is fitting an additive model (greedily).
- The numerical optimization view: boosting makes local improvement iteratively—gradient descent in the function space.
- Gradient boosting is a generic framework
  - Any differentiable loss function
  - Classification, regression, ranking, multiclass etc.
  - Scalable, e.g., XGBoost