

What is Machine Learning

Mengye Ren

NYU

September 5, 2023

Machine Learning Problems

We'll start with a few canonical examples.

What is learning?

"The activity or process of gaining knowledge or skill by studying, practicing, being taught, or experiencing something."

Merriam Webster dictionary

What is learning?

"The activity or process of gaining knowledge or skill by studying, practicing, being taught, or experiencing something."

Merriam Webster dictionary

"A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E ."

Tom Mitchell

Applications of machine learning

- For many problems, it's difficult to program the correct behavior by hand
 - recognizing people and objects
 - understanding human speech

Applications of machine learning

- For many problems, it's difficult to program the correct behavior by hand
 - recognizing people and objects
 - understanding human speech
- Machine learning approach: program an algorithm to automatically learn from data, or from experience. Typically our goal is to solve a prediction problem of the format:
 - Given an **input** x ,
 - **Predict** an **output** y .

Applications of machine learning

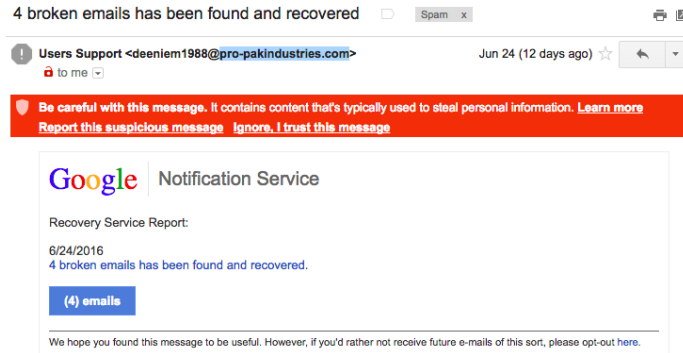
- For many problems, it's difficult to program the correct behavior by hand
 - recognizing people and objects
 - understanding human speech
- Machine learning approach: program an algorithm to automatically learn from data, or from experience. Typically our goal is to solve a prediction problem of the format:
 - Given an **input** x ,
 - **Predict** an **output** y .
- Why might you want to use a learning algorithm?

Applications of machine learning

- For many problems, it's difficult to program the correct behavior by hand
 - recognizing people and objects
 - understanding human speech
- Machine learning approach: program an algorithm to automatically learn from data, or from experience. Typically our goal is to solve a prediction problem of the format:
 - Given an **input** x ,
 - **Predict** an **output** y .
- Why might you want to use a learning algorithm?
 - hard to code up a solution by hand (e.g. vision, speech)
 - system needs to adapt to a changing environment (e.g. spam detection)
 - want the system to perform *better* than the human programmers
 - privacy/fairness (e.g. ranking search results)

Example: Spam Detection

- **Input x:** Incoming email



- **Output y:** "SPAM" or "NOT SPAM"
- This is a **binary classification** problem: there are two possible outputs.

Example: Medical Diagnosis

- **Input x :** Symptoms (fever, cough, fast breathing, shaking, nausea, ...)
- **Output y :** Diagnosis (pneumonia, flu, common cold, bronchitis, ...)
- A **multiclass classification** problem: choosing an output out of a *discrete* set of possible outputs.

How do we express uncertainty about the output?

- **Probabilistic classification** or **soft classification**:

$$\mathbb{P}(\text{pneumonia}) = 0.7$$

$$\mathbb{P}(\text{flu}) = 0.2$$

$$\vdots \quad \quad \vdots$$

Example: Predicting a Stock Price

- **Input x :** History of the stock's prices
- **Output y :** The price of the stock at the close of the next day
- This is called a **regression** problem (for historical reasons): the output is *continuous*.

Comparison to Rule-Based Approaches (Expert Systems)

- Consider the problem of medical diagnosis.
 - ① Talk to experts (in this case, medical doctors).
 - ② Understand how the experts come up with a diagnosis.
 - ③ Implement this process as an algorithm (a **rule-based system**): e.g., a set of symptoms \rightarrow a particular diagnosis.
 - ④ Potentially use logical deduction to infer new rules from the rules that are stored in the knowledge base.

Rule-Based Approach

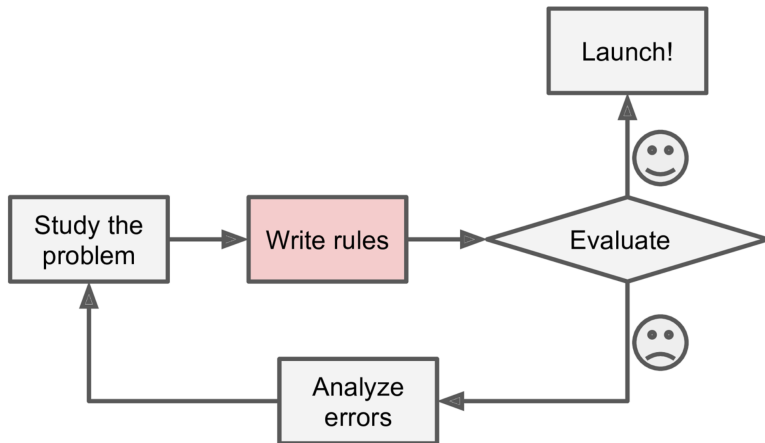


Fig 1-1 from *Hands-On Machine Learning with Scikit-Learn and TensorFlow* by Aurelien Geron (2017).

Advantages of Rule-Based Approaches

- Leverage existing domain expertise.
- Generally **interpretable**: We can describe the rule to another human
- Produce reliable answers for the scenarios that are included in the knowledge bases.

Limitations of Rule-Based Systems

- Labor intensive to build: experts' time is expensive.
- Rules work very well for areas they cover, but often do not **generalize** to unanticipated input combinations.
- Don't naturally handle uncertainty.

The Machine Learning Approach

- Instead of explicitly engineering the process that a human expert would use to make the decision...
- We have the machine **learn** on its own from inputs and outputs (decisions).
- We provide **training data**: many examples of (input x , output y) pairs, e.g.
 - A set of videos, and whether or not each has a cat in it.
 - A set of emails, and whether or not each one should go to the spam folder.
- Learning from training data of this form (inputs and outputs) is called **supervised learning**.

Machine Learning Algorithm

- A **machine learning algorithm** learns from the training data:
 - **Input:** Training Data (e.g., emails x and their labels y)
 - **Output:** A prediction function that produces output y given input x .
- The goal of machine learning is to find the “best” (to be defined) prediction function **automatically, based on the training data**
- The success of ML depends on
 - The availability of large amounts of data;
 - **Generalization** to unseen samples (the test set): just memorizing the training set will not be useful.

Machine Learning Approach

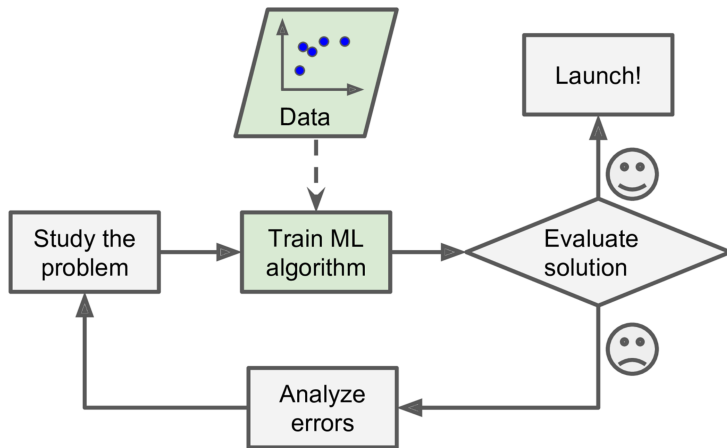


Fig 1-2 from *Hands-On Machine Learning with Scikit-Learn and TensorFlow* by Aurelien Geron (2017).

Key concepts

- The most common **ML problem types**:

Key concepts

- The most common **ML problem types**:
 - Classification (binary and multiclass)

Key concepts

- The most common **ML problem types**:
 - Classification (binary and multiclass)
 - Regression

Key concepts

- The most common **ML problem types**:
 - Classification (binary and multiclass)
 - Regression
- **Prediction function**: predicts output y (e.g. spam or not?) given input x (e.g. email)

Key concepts

- The most common **ML problem types**:
 - Classification (binary and multiclass)
 - Regression
- **Prediction function**: predicts output y (e.g. spam or not?) given input x (e.g. email)
- **Training data**: a set of (input x , output y) pairs

Key concepts

- The most common **ML problem types**:
 - Classification (binary and multiclass)
 - Regression
- **Prediction function**: predicts output y (e.g. spam or not?) given input x (e.g. email)
- **Training data**: a set of (input x , output y) pairs
- **Supervised learning algorithm**: takes training data and produces a prediction function

Key concepts

- The most common **ML problem types**:
 - Classification (binary and multiclass)
 - Regression
- **Prediction function**: predicts output y (e.g. spam or not?) given input x (e.g. email)
- **Training data**: a set of (input x , output y) pairs
- **Supervised learning algorithm**: takes training data and produces a prediction function
- Beyond prediction

Key concepts

- The most common **ML problem types**:
 - Classification (binary and multiclass)
 - Regression
- **Prediction function**: predicts output y (e.g. spam or not?) given input x (e.g. email)
- **Training data**: a set of (input x , output y) pairs
- **Supervised learning algorithm**: takes training data and produces a prediction function
- Beyond prediction
 - **Unsupervised learning**: finding structures in data, e.g. clustering

Key concepts

- The most common **ML problem types**:
 - Classification (binary and multiclass)
 - Regression
- **Prediction function**: predicts output y (e.g. spam or not?) given input x (e.g. email)
- **Training data**: a set of (input x , output y) pairs
- **Supervised learning algorithm**: takes training data and produces a prediction function
- Beyond prediction
 - **Unsupervised learning**: finding structures in data, e.g. clustering
 - **Reinforcement learning**: optimizing long-term objective, e.g. Go

Key concepts

- The most common **ML problem types**:
 - Classification (binary and multiclass)
 - Regression
- **Prediction function**: predicts output y (e.g. spam or not?) given input x (e.g. email)
- **Training data**: a set of (input x , output y) pairs
- **Supervised learning algorithm**: takes training data and produces a prediction function
- Beyond prediction
 - **Unsupervised learning**: finding structures in data, e.g. clustering
 - **Reinforcement learning**: optimizing long-term objective, e.g. Go
 - **Representation learning**: learning good features of real-world objects, e.g. text

Core Questions in Machine Learning

Given any task, the following questions need to be answered:

- **Modeling:** What class of prediction functions are we considering?
- **Learning:** How do we learn the “best” prediction function in this class from our training data?
- **Inference:** How do we compute the output of the prediction function for a new input?